

DATABASE MANAGEMENT

SYSTEMS



WEEK 1

- Drawbacks of using file systems to store data:
- ↳ Data redundancy and inconsistency
 - multiple file formats ; duplication of information
 - ↳ Difficulty in accessing data
 - ↳ Data isolation
 - ↳ Data integrity
 - ↳ Atomicity of updates
 - Failures may leave databases in an inconsistent state with partial data updates carried out.
 - ↳ Concurrent access by multiple users
 - ↳ Security problems

	Bookkeeping	Spreadsheet Files
Durability	Physical damage	computer applications
Scalability	difficult to maintain years of data	easier to search, insert and modify
Security	Susceptible to tampering	can be password-protected
Retrieval	time consuming to search	searching and manipulation requires less manpower
Consistency	Prone to human errors	less prone to mistakes

→ Client / server RDBMS

- ↳ SYBASE
- ↳ ORACLE DATABASE
- ↳ SQL SERVER
- ↳ INFORMIX
- ↳ MySQL

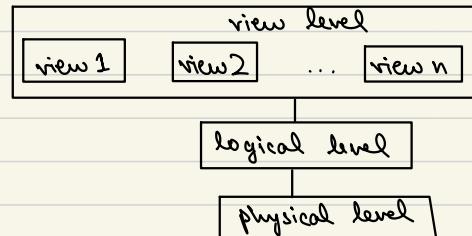
	File handling using Python	DBMS
Scalability (amount of data)	difficult to insert, update and querying of records.	in-built feature to provide high scalability.
Scalability (changes in structure)	extremely difficult to add or remove attributes.	adding/removing attributes seamless using simple SQL queries
Time of execution	in seconds	in milliseconds
Persistence	processing using temp. data structures that have to be manually updated.	automatic, system induced mechanisms
Robustness	has to be done manually	backup, recovery, and restore need minimum manual intervention
Security	difficult to implement in Python	user-specific access at database level.
Programmer's productivity	most operations require extensive code to ensure persistence, robustness, and security	standard and simple in-built queries reduce the effort
Arithmetic operations	easy to do arithmetic computations	limited set of arithmetic operations
Costs	low costs for hardware, software and labour	high costs for hardware, software, labour

→ Levels of Abstraction

- ① Physical level → describes how a record is stored
- ② Logical level → " data stored in database ; relationships among data and fields .

```
type instructor = record
  ID : string;
  name : string;
  dept_name : string;
  salary : integer;
end;
```

- ③ View level → application programs hide details of data types



→ Schema and Instances

→ schema : the way data is organised ; instance : actual data

→ logical schema : analogous to type information of a variable

↳ customer schema :

Name	Customer ID	Account #	Aadhaar ID	Mobile #
------	-------------	-----------	------------	----------

→ physical schema : physical structure

→ instance : actual content ; analogous to value of variable

↳ customer instance

Name	Customer ID	Account #	Aadhaar ID	Mobile #
Pavan Laha	6728	917322	182719289372	9830100291
Lata Kala	8912	827183	918291204829	7189203928
Nand Prabhu	6617	372912	127837291021	8892021892

→ physical data independence : ability to modify the physical schema without changing the logical schema

→ Data Models

- Relational Model → tabular form
- Entity-Relationship data model (for database design)
- Object-based data models (object-oriented and object-relational)

→ DDL (Data Definition Language)

- the way to deal with the schema
- defines schematic of database
- DDL compiler generates a set of table stored in data dictionary

→ DML (Data Manipulation Language) A.R.A. Query language

- for accessing and manipulating data
- Classes:

- ① Pure → Relational Algebra
 - Tuple relational calculus
 - Domain " "
- ② Commercial → SQL

→ SQL → usually embedded in some higher-level language to be able to do complex functions.

→ Database Design

- logical Design → deciding on schema
 - business decision: what attributes?
 - comp. science decision: what relational schema?
- Physical Design → how you physically lay out data / files / etc.

Design Approaches

- ① Entity Relationship Model (Ch. 7)
 - ↳ business requirements
- ② Normalization Theory (Ch. 8)

Object Relational Data Model

- relational model stores atomic values
- object relational data model stores objects

XML • Extensible Markup Language

- ability to specify new tags, and to create nested tag structures
- basis for all new generation data interchange formats

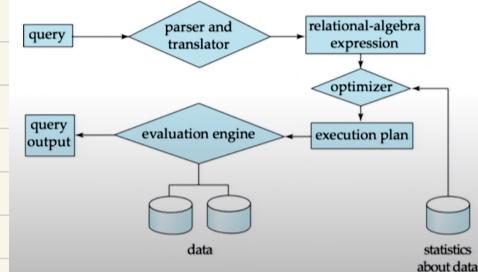
→ Database Engine

→ Storage Manager

- ↳ stores data ; physical design , data dictionary
- ↳ responsible for - interaction with OS file manager
 - efficient storing , retrieving and updating data

→ Query Processing

- ① Parsing and translation
- ② Optimization
- ③ Evaluation



→ Transaction Manager

- transaction : collection of operations that performs in a single logical function in DB application
- transaction - management component : ensures DB remains in a consistent state despite failures
- concurrency - control manager : controls the interaction among concurrent transactions to ensure consistency of DB



WEEK 2

→ Attributes

- normally required to be atomic
- null value is a member of every domain
- Domain of attributes: defines the type . e.g. Name : string
Adhaar ID : 12-digit number

→ Schema and Instance

- A_1, A_2, \dots, A_n are attributes
- $R = (A_1, A_2, \dots, A_n)$ is a relational schema
- D_1, D_2, \dots, D_n are domains of attributes
- a relation is set of n-tuples (a_1, a_2, \dots, a_n) , where $a_i \in D_i$
- Order of tuples / row is irrelevant
No two tuples / row may be identical

→ Key

- let $K \subseteq R$, where R is the set of attributes.
- K is a superkey of R if values of K are sufficient to identify a unique tuple of each possible relation $r(R)$.
- Superkey K is candidate key if K is minimal
 - ↳ no subset of K is a superkey
- Primary key → one of selected candidate key.
- Surrogate key → synthetic key ; not derived from application data.
 - unique identifier for either an entity in modeled world or an object in database.
- Secondary key → non-primary candidate keys
- Simple key → consists of a single attribute
- Composite key → " " multiple "
- Foreign key → when an attribute is a key in a different table

→ Relational Query Languages

① Procedural programming

- ↳ programmer tells the computer what to do
- ↳ HOW to get the output

② Declarative programming

- ↳ more descriptive
- ↳ WHAT relationship holds b/w various entities
- ↳ useful for relational query

→ Select Operation

→ to select a subset of rows from a relational substance that satisfies a condition

→ e.g.

Relation $r =$

A	B	C	D
α	α	1	7
α	β	5	7
β	β	12	3
β	β	23	10

$$\sigma_{A=B \wedge D>5}(r) =$$

A	B	C	D
α	α	1	7
β	β	23	10

↳ $A=B$ AND $D>5$

→ Project Operation

→ to select a few columns

e.g.

A	B	C
α	10	1
α	20	1
β	30	1
β	40	2

$$\pi_{A,C}(r) =$$

A	C
α	1
α	1
β	1
β	2

A	C
α	1
β	1
β	2

→ Union of two relations

A	B
α	1
α	2
β	1

A	B
α	2
β	3

A	B
α	1
α	2
β	1
β	3

→ Difference of two relations

→

A	B
α	1
β	1

→ eliminate those things that are in s that are already in r .

→ Intersection of two relations

→

A	B
α	2

→ Cartesian - product

→

A	B
α	1
β	2

C	D	E
α	10	a
β	10	a
β	20	b
γ	10	b

$$r \times s =$$

A	B	C	D	E
α	1	α	10	a
α	1	β	10	a
α	1	β	20	b
α	1	γ	10	b
β	2	α	10	a
β	2	β	10	a
β	2	β	20	b
β	2	γ	10	b

→ Remaining a table

→ refer to a relation (say E) by more than one name

↳ $p_x(E)$ returns the expression E under the name X .

$r =$

A	B
α	1
β	2

$$r \times p_s(r) =$$

↑

Cartesian with itself

r.A	r.B	s.A	s.B
α	1	α	1
α	1	β	2
β	2	α	1
β	2	β	2

→ Composition of Operations

→ e.g. $\sigma_{A=C}(r \times s)$

$(r \times s) =$

A	B	C	D	E
α	1	α	10	a
α	1	β	10	a
α	1	β	20	b
α	1	γ	10	b
β	2	α	10	a
β	2	β	10	a
β	2	β	20	b
β	2	γ	10	b

$\sigma_{A=C}(r \times s) =$

A	B	C	D	E
α	1	α	10	a
β	2	β	10	a
β	2	β	20	b

→ Natural Join

→ natural join of relations R and S is a relation on schema $R \cup S$:

→ t_r, t_s tuples from r and s

→ if t_r and t_s have same values on each of the attributes $R \cap S$ (common attributes), add a tuple t to result, where:

t has same value as t_r on r
 t " " " " " t_s " s

→ example:

Relations =

A	B	C	D	
α	1	α	a	
β	2	γ	a	
γ	4	β	b	
α	1	γ	a	
δ	2	β	b	

B	D	E
1	a	α
3	a	β
1	a	γ
2	b	δ
3	b	ε

$r \bowtie s =$

A	B	C	D	E
α	1	α	a	α
α	1	α	a	γ
α	1	γ	a	α
α	1	γ	a	γ
δ	2	β	b	δ

→ $\pi_{A, r.B, C, r.D, E}(\sigma_{r.B=s.B \wedge r.D=s.D}(r \times s))$

→ Aggregate Operators

→ SUM, AVG, MAX, MIN

→ SQL

History → developed by IBM
 → pronounced SEQUEL

Data Definition Language (DDL)

- Schema for each relation
- Domain of values of each attribute
- Integrity constraints

→ relation defined using `create table` command:

`create table r (A1, D1, A2, D2, ..., An, Dn),`
 `(integrity-constraint1),`
 `...`
 `(integrity-constraintk);`

attribute name in schema of relation*

data type of values in the domain of the attribute

example :

- not null ✓
- primary key (A₁, ..., A_n) ✓
- foreign key (A_m, ..., A_n) references r ✓

```
create table instructor (  
    ID char(5),  
    name varchar(20)  
    dept_name varchar(20)  
    salary numeric(8, 2));
```

```
create table instructor (  
    ID char(5),  
    name varchar(20) not null,  
    dept_name varchar(20),  
    salary numeric(8, 2),  
    primary key (ID),  
    foreign key (dept_name) references department);
```

primary key declaration on an attribute automatically ensures not null

→ Update tables

→ Insert (DML) : add a value

→ Delete (DML) : remove all rows from table

→ Drop Table (DDL) : discard the entire table

→ Alter (DDL) : add attribute or drop attribute

Data Manipulation Language (DML)

→ Typical SQL Query:

```
select attributes  
from relations  
where condition (Boolean)
```

→ Select → corresponds to projection operation of relational algebra

→ SQL names are case-insensitive

→ allows duplicates in relations and query results

 ↳ use `distinct` to remove duplicates

→ `select *` → selects all attributes

→ can contain arithmetic operations (+, -, *, /)

→ Where → condition, True or False

→ corresponds to selection operation of relational algebra

→ From → scope of which relations you're dealing with

→ for multiple relations, by default, uses Cartesian product

→ String Operation

→ percent (%) → character matches any substring

→ underscore (_) → character matches any character

→ Ordering display of tuples ; e.g. → `select distinct name
from instructor
order by name desc`

→ Select top clause ; e.g. → select top 10 distinct name
from instructor

→ Set operations

→ union, intersect, except

→ example .

- Find the salaries of all instructors that are less than the largest salary
select distinct T.salary
from instructor as T, instructor as S
where T.salary < S.salary

} will give highest salary

- Find all the salaries of all instructors
select distinct salary
from instructor

- Find the largest salary of all instructors
(select "second query")
except
(select "first query")

} will give all salaries except the highest salary

→ to retain all duplicates: union all, except all, intersect all

→ Null values

- usually for when you don't know all the values of all the fields .
- adding an attribute to a table .
- the condition where is null can be used to check for null values . where is not null
- any comparison with null returns unknown
- three-valued logic

OR : (unknown or true) = true

(unknown or false) = unknown

(unknown or unknown) = unknown

AND : (unknown and true) = unknown

(unknown and false) = false

(unknown and unknown) = unknown

NOT : (not unknown) = unknown

→ (P is unknown) = true if P evaluates to unknown

→ if the result of condition is unknown , it is treated as false in the "where" statement .

Binary Numbers and Character Encoding

→ To convert a number into equivalent binary representation : divide the number by 2 until the quotient becomes 0 . The remainder after every division (in the order) is the binary representation .

e.g. $(1729)_{10} \Rightarrow (11011000000)_2$

$$\begin{array}{r} 2 | 1729 \\ 2 | 864 - 1 \\ 2 | 432 - 0 \\ \vdots \\ 2 | 13 - 1 \\ 2 | 6 - 1 \\ 2 | 3 - 0 \\ 2 | 1 - 1 \\ 0 - 1 \end{array}$$

$(2.718)_{10} \approx (10 \cdot 10110)_2$

$$\begin{array}{r} 2 | 2 \\ 2 | 1 - 0 \\ 2 | 0 - 1 \end{array}$$

$$\begin{array}{r} 0.718 \times 2 = 1.436 \\ 0.436 \times 2 = 0.872 \\ 0.872 \times 2 = 1.744 \\ 0.744 \times 2 = 1.488 \\ 0.488 \times 2 = 0.976 \end{array}$$

$(0.718)_{10} = (10110)_2 \rightarrow$ not as precise because we didn't get to 0 .

→ To convert from binary to number: $\sum_{i=0}^{n-1} x_i \cdot (2)^i$, where $n = \text{total # of digits}$, and i counts backward starting from the rightmost, and $x_i = i\text{th digit}$.

e.g. $11011000001 \rightarrow^*$

$$1 \cdot (2)^0 + 1 \cdot (2)^1 + 0 \cdot (2)^2 + \dots + \\ 0 \cdot (2)^5 + 0 \cdot (2)^6 + 1 \cdot (2)^7 = 1729$$



WEEK 3

→ Nested Subqueries

→ 'select - from - where' nested within another query.

→ As a part of WHERE clause:

→ typical use → set membership / comparison / cardinality

Set Membership

example :

- Find courses offered in Fall 2009 and in Spring 2010. (intersect example)
- ```
select distinct course_id
from section
where semester = 'Fall' and year = 2009 and
course_id in (select course_id
from section
where semester = 'Spring' and year = 2010);
```

↳ nested  
subquery

→ SOME clause

→  $F < \text{comp} > \text{some } r \Leftrightarrow \exists t \in r \text{ such that } (F < \text{comp} > t)$

where  $< \text{comp} >$  can be :  $<, \leq, >, \geq, =, \neq$

→ existential quantification

example :

#### Set Comparison – "some" Clause

- Find names of instructors with salary greater than that of some (at least one) instructor in the Biology department

```
select distinct T.name
from instructor as T, instructor as S
where T.salary > S.salary and S.dept_name = 'Biology';
```

- Same query using some clause

```
select name
from instructor
where salary > some (select salary
from instructor
where dept_name = 'Biology');
```



→ ALL clause

→  $F < \text{comp} > \text{all } r \Leftrightarrow \forall t \in r \text{ such that } (F < \text{comp} > t)$

↳ comparison is true for all.

→ universal quantification

→ EXISTS clause

→ returns true if argument subquery is nonempty

#### Use of "exists" Clause

- Yet another way of specifying the query "Find all courses taught in both the Fall 2009 semester and in the Spring 2010 semester"

```
select course_id
from section as S
where semester = 'Fall' and year = 2009 and
exists (select *
from section as T
where semester = 'Spring' and year = 2010
and S.course_id = T.course_id);
```

- Correlation name – variable  $S$  in the outer query

- Correlated subquery – the inner query



## → UNIQUE clause

→ tests whether a subquery has any duplicates

example:

Find all courses that were offered at most once in 2009

```
select T.course_id
from course as T
where unique (select R.course_id
from section as R
where T.course_id = R.course_id
and R.year = 2009);
```

## → As a part of FROM clause

→ Find the average instructors' salaries of those departments where the average salary is greater than \$42,000

```
select dept_name, avg_salary
from (select dept_name, avg(salary) as avg_salary
 from instructor
 group by dept_name)
 where avg_salary > 42000;
```

## → WITH clause

→ defining a temporary relation whose definition is only available to the query in which the with clause is.

example: creates a table with one attribute

Find all departments with the maximum budget

```
with max_budget(value) as
 (select max(budget)
 from department)
select department.name
 {
 from department, max_budget
 where department.budget=max_budget.value;
 } → basic query
```

## Complex Queries using With Clause

example:

- Find all departments where the total salary is greater than the average of the total salary at all departments

```
with dept_total (dept_name, value) as
 select dept_name, sum(salary)
 from instructor
 group by dept_name,
 dept_total.avg(value) as
 (select avg(value)
 from dept_total)
select dept_name
 from dept_total, dept_total_avg
 where dept_total.value > dept_total_avg.value;
```

## → As a part of SELECT clause

→ needs scalar subquery

example :

List all departments along with the number of instructors in each department

```
select dept_name,
 (select count(*)
 from instructor
 where department.dept_name = instructor.dept_name)
 as num_instructors
 from department;
```

→ Modifications of the Database

→ Deletion of a tuple

examples:

→ Delete all tuples: delete from instructors

→ Delete with condition: delete from instructors  
where dept-name = 'Finance'

→ Insertion of a tuple

examples:

→ Add a new tuple: insert into course  
values ('CS-437', ..., 4)

→ Updation of a tuple

Increase salaries of instructors whose salary is over \$100,000 by 3%, and all others by a 5%

- o Write two update statements:

```
update instructor
set salary = salary * 1.03
where salary > 100000;
update instructor
set salary = salary * 1.05
where salary <= 100000;
```

CASE statement for conditional updates



- Same query as before but with case statement

```
update instructor
set salary = case
when salary <= 100000
then salary * 1.05
else salary * 1.03
end
```

→ CROSS JOIN

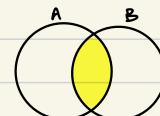
→ returns cartesian product of rows

select \*

from employee, department

→ INNER JOIN

→ where only the intersecting tuples are taken

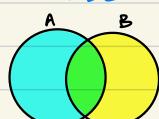
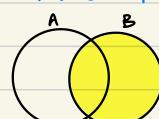
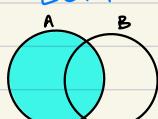


→ ignores uncommon tuples

→ natural join removes duplicity of common column

→ OUTER JOIN

→ LEFT or RIGHT or FULL



natural left/right/full outer join

## → VIEW

→ defined using `create view` statement:

`create view v as <query expression>`  
creates a virtual relation  
with the name 'v'.

→ examples:

- A view of instructors without their salary  
`create view faculty as  
select ID, name, dept_name  
from instructor`
- Find all instructors in the Biology department  
`select name  
from faculty  
where dept_name = 'Biology'`
- Create a view of department salary totals  
`create view departments_total_salary(dept_name, total_salary) as  
select dept_name, sum(salary)  
from instructor  
group by dept_name;`

→ view recursive: a view relation that depends on itself to be defined.

→ insertion in a view is actually done by insertion in the relation.

## → Transactions

→ unit of work ; atomic (fully executed or entirely rolled back)

→ ended by `commit work` or `rollback work`

## → Integrity Constraints in a single relation

→ `not null`

↳ declare an attribute to not allow null values

→ `unique (A1, A2, ..., Am)`

↳ the attributes form a candidate key.

→ `check`

↳ to impose business logic

↳ e.g. `check (semester in ('Fall', 'Winter', 'Spring'))`

## → Referential Integrity

→ a value that appears in one relation for a given set of attributes also appears for a certain set of attributes in another relation.

→ example:

`foreign key (dept_name) references department`

{  
  `on delete cascade`  
  `on update cascade`

↳ maintains

integrity; if key deleted in foreign relation, then all tuples with that key will be deleted in this relation.

## → SQL Datatypes and Schemas

→ date, time, timestamp, interval

## → Index creation

→ create index index-name on relation(attribute)

## → User-Defined Types (UDT)

→ create type Dollars as numeric(12,2) final

## → Domains

→ create domain person-name char(20) not null

→ create domain degree-level varchar(10)

constraint degree-level-test

check (value in ('Bachelors', 'Masters', 'Doctorate'));

## → Authorisation

→ grant <privilege list>  
on <relation/view name> to <user list>

### → Privileges

→ select : read access

→ insert : ability to insert tuples

→ update : ability to use SQL update statement

→ delete : delete tuples

→ all privileges : for all allowable privileges

→ revoke → revokes the privileges

### → Roles

↳ creates a role ; then you grant privileges to roles and roles can be granted to users.

↳ create role <role-name>

↳ roles can be granted to other roles, creating a hierarchy of authorisation.

→ Function and Procedures  
→ example:

### SQL Functions

- Define a function that, given the name of a department, returns the count of the number of instructors in that department:

```
create function dept_count (dept_name varchar(20))
returns integer
begin
declare d_count integer;
select count(*) into d_count
from instructor
where instructor.dept_name = dept_name
return d_count;
end
```

- The function `dept_count` can be used to find the department names and count of all departments with more than 12 instructors:

```
select dept_name, budget
from department
where dept_count(dept_name) > 12
```



→ function is a parameterised view

→ Table functions

↳ returns a table

↳ example:

• Functions that return a relation as a result added in SQL:2003

• Return all instructors in a given department:

```
create function instructor_of (dept_name char(20))
returns table (
ID varchar(5),
name varchar(20),
dept_name varchar(20),
salary numeric(8,2))
returns table
(select ID, name, dept_name, salary
from instructor
where instructor.dept_name = instructor_of.dept_name)
```

• Usage

```
select *
from table (instructor_of ('Music'))
```

→ Procedures

↳ do not return anything. They have 'in' parameters and 'out' parameters.

↳ example:

• The `dept_count` function could instead be written as procedure:

```
create procedure dept_count_proc (
in dept_name varchar(20), out d_count integer)
begin
select count(*) into d_count
from instructor
where instructor.dept_name = dept_count_proc.dept_name
end
```

• Procedures can be invoked either from an SQL procedure or from embedded SQL, using the `call` statement:

to execute a procedure

```
declare d_count integer;
call dept_count_proc('Physics', d_count);
```

• Procedures and functions can be invoked also from dynamic SQL

• SQL:1999 allows **overloading** - more than one function/procedure of the same name as long as the number of arguments and / or the types of the arguments differ

→ Language Constructs

→ begin ... end → compound statements

→ while boolean expression do  
sequence of statements;  
end while;

→ repeat  
sequence of statements;  
until boolean expression  
end repeat;

→ Find the budget of all departments:

```
declare n integer default 0;
for r as
 select budget from department
 do
 set n = n + r.budget
 end for;
```

→ if - then - else

→ case variable

```
when value1 then
 sequence of statements
when value2 then
 sequence of statements
...
else
 sequence of statements
end case;
```

→ Searched case example:

#### Language Constructs (6): Searched case

- Searched case statement
 

```
case
 when sql-expression = value1 then
 sequence of statements;
 when sql-expression = value2 then
 sequence of statements;
 ...
 else
 sequence of statements;
 end case;
```

## Triggers

→ set of actions performed in response to an **insert**, **delete**, or **update** operation.

→ **create trigger** statement

### Statement level triggers

↳ applies to all rows of the statement result

### Row level triggers

↳ applies to a single row

→ example:

#### Trigger to Maintain credits\_earned value

```
create trigger credits_earned after update of grade on takes
referencing new row as nrow
referencing old row as orow
for each row
when nrow.grade <> 'F' and nrow.grade is not null
 and (orow.grade = 'F' or orow.grade is null)
begin atomic
 update student
 set tot_cred= tot_cred +
 (select credits
 from course
 where course.course_id=nrow.course_id)
 where student.id = nrow.id;
end;
```

# WEEK 4

→ Formal relational query language systems:

- ① Relational algebra → procedural and algebra based
- ② Tuple relational calculus → non-procedural and predicate calculus
- ③ Domain relational calculus → " " " "

→ Relational Algebra

→ basic operators:

|             |                       |         |
|-------------|-----------------------|---------|
| select : σ  | set difference : -    | and : ∧ |
| project : π | cartesian product : × | or : ∨  |
| union : ∪   | rename : ρ            | not : ¬ |

→ Division

→ example:

| R        |           | S         | R   S    |
|----------|-----------|-----------|----------|
| Lecturer | Module    | Subject   | Lecturer |
| Brown    | Compiler  | Databases | Green    |
| ✓ Brown  | Databases | Prolog    |          |
| ✓ Green  | Prolog    |           |          |
| ✓ Green  | Databases |           |          |
| ✓ Lewis  | Prolog    |           |          |
| ✓ Smith  | Databases |           |          |

'Green' is associated with all instances in S relation.

↳ needs one common attribute.

↳ looks for (R-S) attributes that have the same values on common attribute.

↳ in the example, 'Brown' Lecturer didn't have an entry for 'Prolog' Modules so disqualified.

'Green' Lecturer is the only value that has the same value for both 'Prolog' and 'Databases'.

→ example:

| A  | B1  | A/B1 | sno |
|----|-----|------|-----|
|    | pno |      | s1  |
| s1 | p1  |      | s2  |
| s1 | p2  |      | s3  |
| s1 | p3  |      | s4  |
| s1 | p4  |      |     |
| s2 | p1  |      |     |
| s2 | p2  |      |     |
| s3 | p2  |      |     |
| s4 | p2  |      |     |
| s4 | p4  |      |     |

| B2 | B3  | A/B2 | sno |
|----|-----|------|-----|
|    | pno |      | s1  |
|    | p2  |      | s2  |
|    | p4  |      | s3  |

| A/B3 | sno |
|------|-----|
|      | s1  |

## → Predicate Logic

→ extension of propositional logic / boolean algebra

→ Predicate is a function.

↳ example:  $P(n) = n > 3$

↳ once you assign a value to 'n', it becomes a proposition and has a truth/false value.

↳ can have multiple variables:  $P(n_1, n_2, n_3, \dots, n_n)$

→ Quantifiers → used to talk about properties that hold over the entire domain.

↳ Universal quantifiers → has to hold true for all elements in the domain.

↳ example:  $\forall P(n)$  read as "for all  $n$   $P(n)$ "

↳  $P(n) = n + 2 > n$ ; then  $\forall n P(n)$  is True

↳ Existential quantifiers → has to hold for at least one element

↳ example:  $\exists n P(n)$

↳  $P(n) = n > 5$ ; then  $\exists n P(n)$  is True

## → Tuple Relational Calculus

→  $\{t \mid P(t)\}$

→ example: find first names of all students with age  $> 21$ .

↳  $\{t.Fname \mid t \in \text{student}(t) \wedge t.age > 21\}$

## → Domain Relational Calculus

→ no diff. from tuple relational other than notation.

## → Equivalence of RA, TRC, DRC

→ example: Select operation on a relation  $R = (A, B)$

RA:  $\sigma_{B=17}(r)$       TRC:  $\{t \mid t \in r \wedge t.B = 17\}$

DRC:  $\{\langle A, B \rangle \mid \langle A, B \rangle \in r \wedge B = 17\}$

## → Design of an E-R model

→ satisfies functional specification

→ conforms to limitations

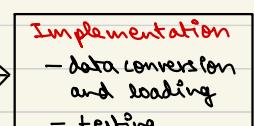
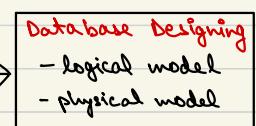
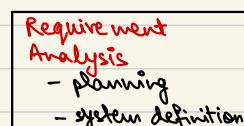
→ meets requirements on performance and resource usage

→ satisfies design criteria

→ " restrictions on the design process

→ Two components: ① Abstraction

② Model



## → Abstraction of a design

→ abstract things down so we have the information but it needs a lot less memory | time | effort.

→ example: when dealing with courses / grades, salary of the instructor is irrelevant.

## → Models of a design

→ models of the real world that mimic the real world.

## → Entity - Relationship Model

→ Framework for logical model

## → Conceptual Requirements

## → Attributes

↳ properties associated with an entity / entity set

→ simple and composite < <sup>first name</sup>  
~~last name~~

↳ single-valued or multivalued

↳ example : multiple cars of a single person

↳ derived attributes

→ example: age, given DOB

↳ domain → set of permitted values

→ Entity sets

↳ represented by a list of attributes

↳ relations | tables

→ Relationship sets

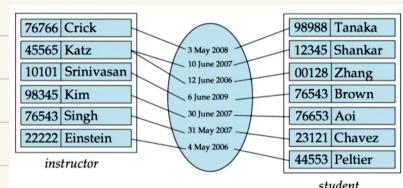
↳ association among several entities.

→ example : 44553 (Peltier) student entity advisor relationship set 22222 (Einstein) instructor entity

$\hookrightarrow (44553, 22222) \in \text{advisor}$

↳ relationship could itself have an attribute

↳ example :



## ↳ Mapping cardinalities

→ One-to-one

→ One - to - Many

→ Many - to - One

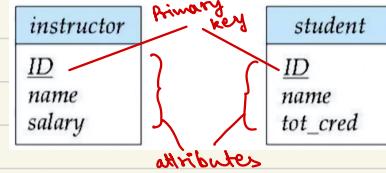
→ Many - to - Many

→ Weak Entity Set

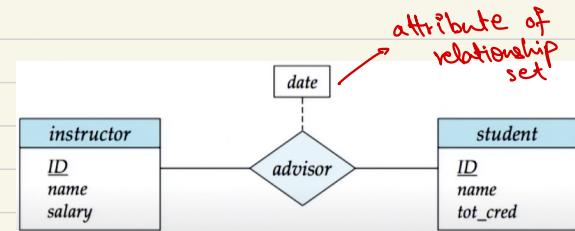
- strong entity sets have a primary key.
- weak " " have a partial key, set of (discriminator) attributes that uniquely identify a group of entities.
- They feature in relationship with a strong entity set

→ ER diagram

→ Entity Sets  
↳ Box

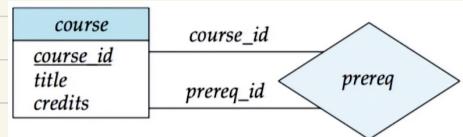


→ Relationship Sets  
↳ Diamond



→ Roles

↳ entity set may have relationship with itself.



→ Cardinality constraints

↳ one-to-many and so on.

↳ directed line (→) means "one".

undirected " (—) means "many".

→ Total / Partial participation

↳ Total : (==) every tuple participates in the relationship.

↳ Partial : (—) not every tuple participates.

→ Complex cardinality

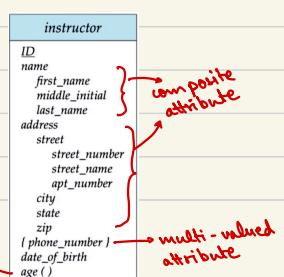
↳ min/max cardinality shown in the form:

$l \dots h$ , where  $l = \min$  and  $h = \max$

→ Complex attributes

↳ shown with a tail.

example.



→ Weak entity sets

↳ double-lined boundary

↳ diamond (relationship) will also be double-lined.

↳ discriminator underlined with dashed.



→ ER model extended features

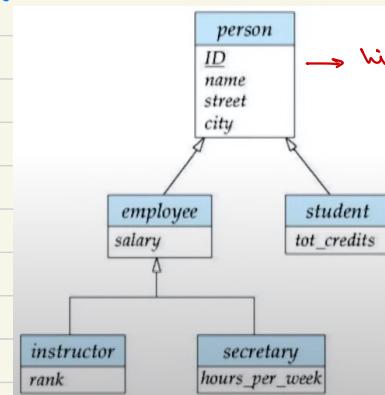
→ Specialisation → top-to-bottom design

→ ISA

↳ if  $r_1$  ISA  $r_2$ , then  $r_1$  has all the attributes of  $r_2$ . Depicted by ( $\rightarrow$ )

Overlapping: employee and person

Disjoint: instructor and secretary



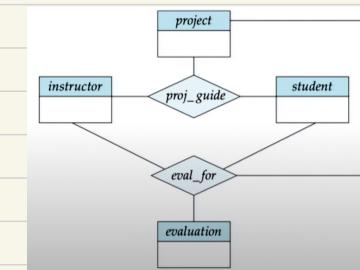
→ highest entity set

- If it is not necessary that every entity in a higher entity set exists in one or more lower entity sets, then the specialisation is partial otherwise total.

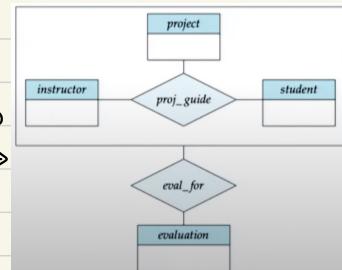
→ Generalisation → bottom-to-top design

↳ tends to be total but not always the case.

→ Aggregation



can be converted into



→ proj\_guide can be treated as an entity, to remove redundancy.

→ eval\_for is eval\_for(s\_ID, project\_id, i\_ID, evaluation\_id)



# WEEK 5

- ## → Features of a good relational design

- reflects real-world structure
  - can represent all expected data over time
  - avoids redundant storage of data
  - provides efficient access to data
  - supports maintenance of data integrity
  - clean, consistent, and easy to understand

- Redundancy → multiple copies of same data ; leads to anomaly

- Anomaly → inconsistencies due to changes in database

- Ingestion anomaly

↳ happens when insertion of a data record is not possible without adding some additional unrelated data.

- Deletion anomaly

- ↳ when deletion of a record results in losing some unrelated information

- Update anomaly

- ↳ when a data is changed, many records have to be changed leading to possibility of some changes being made incorrectly.

- To minimise redundancy, minimise dependency

- to "dependency, you need a good decomposition of the database.

- ## → Lossless Join Decomposition

- a decomposition of a relation  $R$  into relations  $R_1, R_2$  such that if we perform natural join of  $R_1$  and  $R_2$ , it will return  $R$ .

- ## → Atomic Domains and First Normal Form

- a domain is atomic if its elements are indivisible units.  
↳ example of non-atomic :- set of names  
- id numbers like CS101

- a relational schema R is in **First Normal Form (1NF)** if
    - domains of all attributes are **atomic**.
    - every attribute contains only a **single-value**.

## → Functional Dependencies

- value of certain attributes determining uniquely the value for another set of attributes.
- $\alpha \subseteq R$  and  $\beta \subseteq R$   
 $\alpha \rightarrow \beta$  holds on  $R$  iff  $\forall$  legal relation  $r(R)$ , for any two tuples  $t_1$  and  $t_2$ ; if  $t_1[\alpha] = t_2[\alpha] \Rightarrow t_1[\beta] = t_2[\beta]$
- $K$  is a **superkey** for relational schema  $R$  iff  $K \rightarrow R$   
 $K$  is a **candidate key** for  $R$  iff  
 $K \rightarrow R$ ; and  
for no  $\alpha \subset K$ ,  $\alpha \rightarrow R$

- **Reflexivity**: if  $\beta \subseteq \alpha$ , then  $\alpha \rightarrow \beta$
- **Augmentation**: if  $\alpha \rightarrow \beta$ , then  $\gamma\alpha \rightarrow \gamma\beta$
- **Transitivity**: if  $\alpha \rightarrow \beta$  and  $\beta \rightarrow \gamma$ , then  $\alpha \rightarrow \gamma$
- Axioms can be applied to generate new FDs
  - ↳ a new FD obtained by applying the axioms is **logically implied**
  - ↳ Axioms are:

**Sound**: generate only FDs that actually hold  
**Complete**: eventually generate all FDs that hold

- Closure set of FDs  $F = F^+$ 
  - ↳ set of all FDs logically implied by  $F$
- **Union**: if  $\alpha \rightarrow \beta$  and  $\alpha \rightarrow \gamma$ , then  $\alpha \rightarrow \beta\gamma$
- **Decomposition**: if  $\alpha \rightarrow \beta\gamma$ , then  $\alpha \rightarrow \beta$  and  $\alpha \rightarrow \gamma$
- **Pseudotransitivity**: if  $\alpha \rightarrow \beta$  and  $\gamma\beta \rightarrow \delta$ , then  $\alpha\gamma \rightarrow \delta$

} derived rules

## → Closure of attribute sets

- $R = (A, B, C, G_1, H, I)$ ;  $F = \{A \rightarrow B, A \rightarrow C, CG_1 \rightarrow H, CG_1 \rightarrow I, B \rightarrow H\}$
- $(AG_1)^+$  steps:

①  $AG_1$

②  $ABC G_1$  because  $(A \rightarrow C)$  and  $(A \rightarrow B)$

③  $ABC G_1 H$  because  $(CG_1 \rightarrow H)$  and  $(G_1 \subseteq ABC G_1)$

④  $ABC G_1 H I$  because  $(G_1 \rightarrow I)$  and  $(G_1 \subseteq ABC G_1 H)$

→  $(AG_1)^+$  is a super key because  $(AG_1)^+$  has all attributes of  $R$ .

→  $(AG_1)^+$  is minimal because  $(A)^+ = (ABCH)$  and  $(G_1)^+ = (G_1)$   
hence  $(AG_1)^+$  is a candidate key.

## → Decomposition using Functional Dependencies

### → BCNF: Boyce - Codd Normal Form

→ A relational schema  $R$  is in BCNF w.r.t. a set  $F$  of FDs  
if for all FDs in  $F^+$  satisfy one of these:

- ①  $\alpha \rightarrow \beta$  is trivial (i.e.,  $\beta \subseteq \alpha$ ); or
- ②  $\alpha$  is a superkey for  $R$

→ **Decomposition**: if in R a non-trivial FD  $\alpha \rightarrow \beta$  violates BCNF, we decompose R into:

- ①  $\alpha \cup \beta$
- ②  $R - (\beta - \alpha)$

→ It is not always possible to achieve both BCNF and dependency preservation.

→ **Third Normal Form (3NF)**

→ R is in 3NF if for all  $\alpha \rightarrow \beta \in F^+$  at least one of the following holds:

- ①  $\alpha \rightarrow \beta$  is trivial (i.e.,  $\beta \subseteq \alpha$ ) ; or
- ②  $\alpha$  is a superkey for R ; or
- ③ each attribute A in  $(\beta - \alpha)$  is contained in a contained key for R.

→ is a relaxation of BCNF to preserve dependencies.

→ **Algorithms for Functional Dependencies**

→ **Attribute set Closure**

→ Use cases: testing for superkey / candidate key; testing for functional dependencies; and computing closure of F

→ **Extraneous Attributes** (consider FD  $\alpha \rightarrow \beta$ )

→ Attribute A is **extraneous** in  $\alpha$  if  $A \in \alpha$  and  $F$  logically implies  $(F - \{\alpha \rightarrow \beta\}) \cup \{(\alpha - A) \rightarrow \beta\}$ .

→ Attribute A is **extraneous** in  $\beta$  if  $A \in \beta$  and  $(F - \{\alpha \rightarrow \beta\}) \cup \{ \alpha \rightarrow (\beta - A) \}$  logically implies F.

→ **To test if A is extraneous in  $\alpha$ :**

- ① Compute  $(\{\alpha\} - A)^+$  w.r.t. F
- ② Check if  $\beta$  is in  $(\{\alpha\} - A)^+$

→ **To test if A is extraneous in  $\beta$ :**

- ① Compute  $(\alpha)^+$  w.r.t.  $F'$ ,  
where  $F' = (F - \{\alpha \rightarrow \beta\}) \cup (\alpha \rightarrow \{\beta - A\})$
- ② Check if  $(\alpha)^+$  contains A

→ **Equivalence of sets of Functional Dependencies**

→ F & G are equivalent if  $F^+ = G^+$ .

→ Canonical Cover

→  $F_c$  is a canonical cover for  $F$  if:

①  $F^+ = F_c^+$

② no FD in  $F_c$  contains extraneous attribute

③ each left side of FD in  $F_c$  is unique; i.e., there are no two  $\alpha_1 \rightarrow \beta_1$  and  $\alpha_2 \rightarrow \beta_2$  in such that  $\alpha_1 \rightarrow \alpha_2$

→ is a minimal set of FDs

→ Algorithm: To compute a canonical cover for  $F$ :

repeat

    Use the union rule to replace any dependencies in  $F$

$\alpha_1 \rightarrow \beta_1$  and  $\alpha_1 \rightarrow \beta_2$  with  $\alpha_1 \rightarrow \beta_1\beta_2$

    Find a functional dependency  $\alpha \rightarrow \beta$  with an extraneous attribute either in  $\alpha$  or in  $\beta$

    /\* Note: test for extraneous attributes done using  $F_c$ , not  $F$  \*/

    If an extraneous attribute is found, delete it from  $\alpha \rightarrow \beta$

until  $F$  does not change

→ Lossless Join Decomposition

→ Conditions: ①  $R_1 \cup R_2 = R$

②  $R_1 \cap R_2 \neq \emptyset$

③  $R_1 \cap R_2 \rightarrow R_1$  or  $R_1 \cap R_2 \rightarrow R_2$

→ Dependency Preservation

→  $F_i$  is the set of dependencies in  $F^+$  that include only attributes in  $R_i$   
decomposition is dependency preserving if

$$(F_1 \cup F_2 \cup \dots \cup F_n)^+ = F^+$$

→ Superkey, candidate key, prime attributes

→ Superkey determines every tuple. Any attribute 'A' that does not appear on the R.H.S. of any FD must be in every superkey.

→ Candidate key → determines every tuple AND no subset of attributes of candidate key should form superkey.

→ Prime attributes → attributes that are a part of at least one candidate key



# WEEK 6

## → Desirable Properties of Decomposition

- ① Lossless Join Decomposition
- ② Dependency Preserving

→ 3NF relations : free of insertion, update and deletion anomalies.

## → 1NF : First Normal Form

- iff all domains contain atomic values only
- example:

| Students |       |        |
|----------|-------|--------|
| SID      | Sname | Cname  |
| S1       | A     | C,C++  |
| S2       | B     | C++,DB |
| S3       | A     | DB     |

SID : Primary Key  
MVA exists  $\Rightarrow$  Not in 1NF

| Students |       |       |
|----------|-------|-------|
| SID      | Sname | Cname |
| S1       | A     | C     |
| S1       | A     | C++   |
| S2       | B     | C++   |
| S2       | B     | DB    |
| S3       | A     | DB    |

SID, Cname : Primary Key  
No MVA  $\Rightarrow$  In 1NF

→ cause of redundancy : if  $X \rightarrow Y$  is a non-trivial FD over R with X is not a superkey of R, then redundancy b/w X and Y.

## → 2NF : Second Normal Form

- must be in 1NF and no partial dependency
- partial dependency :

in  $R(X, Y, A)$ , IF  $Y \rightarrow A$ , where

Y : proper subset of candidate key, and

A : non-prime attribute. Then, R not in 2NF

→ example :

STUDENT(Sid, Sname, Cname) (already in 1NF)

| SID | Sname | Cname |
|-----|-------|-------|
| S1  | A     | C     |
| S1  | A     | C++   |
| S2  | B     | C++   |
| S2  | B     | DB    |
| S3  | A     | DB    |

(SID, Cname) : Primary Key

Redundancy?  
o Sname  
Anomaly?  
o Yes

Functional Dependencies:  
 $\{SID, Cname\} \rightarrow Sname$   
 $SID \rightarrow Sname$

Partial Dependencies:  
 $SID \rightarrow Sname$  (as SID is a Proper Subset of Candidate Key  $\{SID, Cname\}$ )

Key Normalization

| R1: |       | R2: |       |
|-----|-------|-----|-------|
| SID | Sname | SID | Cname |
| S1  | A     | S1  | C     |
| S2  | B     | S1  | C++   |
| S3  | A     | S2  | C++   |

(SID) : Primary Key

(SID, Cname) : Primary Key

The above two relations R1 and R2 are  
1.Lossless Join  
2.2NF  
3.Dependency Preserving

## → 3NF : Third Normal Form

→ must be in 2NF and for every FD  $X \rightarrow A$  either :

- ①  $A \subseteq X$  (i.e., trivial FD) ; or
- ②  $X$  is a superkey ; or
- ③  $A$  is a prime attribute.

→ **transitive dependency** :

①  $A \rightarrow B$ , and ② not  $B \rightarrow A$ , and ③  $B \rightarrow C$ , then  
 $A \rightarrow C$  is a transitive dependency.

→ **remove transitive dependency** to get 3NF.

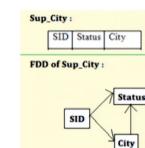
→ **example** :

Sup.City(SID, Status, City) (already in 2NF)

| Sup_City: |        |        |
|-----------|--------|--------|
| SID       | Status | City   |
| S1        | 30     | Delhi  |
| S2        | 10     | Karnal |
| S3        | 40     | Rohtak |
| S4        | 30     | Delhi  |

SID: Primary Key

- Redundancy?
- Status
- Anomaly?
- Yes



Functional Dependencies:

$SID \rightarrow Status$ ,  
 $SID \rightarrow City$ ,  
 $City \rightarrow Status$   
**Transitive Dependency :**  
 $SID \rightarrow Status$   
{As SID → City and City → Status}

| Post Normalization |                   |
|--------------------|-------------------|
| SC:                | CS:               |
| SID                | City              |
| S1                 | Delhi             |
| S2                 | Karnal            |
| S3                 | Rohtak            |
| S4                 | Delhi             |
| SID: Primary Key   | City: Primary Key |

The above two relations SC and CS are  

- Lossless Join
- 3NF
- Dependency Preserving

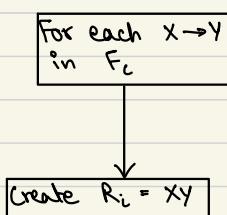
## → Decomposition into 3NF

→ Testing for 3NF

→ no need to check on  $F^+$

→ possible in **polynomial time**.

→ **Algorithm** :



→ **example** :

- ① Compute canonical cover  $F_c$  of  $F$ .
- ② Create a relation  $R_i = XY$  for each FD  $X \rightarrow Y$  in  $F_c$ .
- ③ If key  $K$  of  $R$  does not occur in any relation  $R_i$ , create one more relation  $R_i = K$ .

- Relation schema:  
 $cust\_banker\_branch = (customer\_id, employee\_id, branch\_name, type)$

- The functional dependencies for this relation schema are:

- a)  $customer\_id, employee\_id \rightarrow branch\_name, type$
- b)  $employee\_id \rightarrow branch\_name$
- c)  $customer\_id, branch\_name \rightarrow employee\_id$

- We first compute a canonical cover

- $branch\_name$  is extraneous in the RHS of the 1<sup>st</sup> dependency
- No other attribute is extraneous, so we get  $F_c =$   
 $customer\_id, employee\_id \rightarrow type$   
 $employee\_id \rightarrow branch\_name$   
 $customer\_id, branch\_name \rightarrow employee\_id$

- The for loop generates following 3NF schema:

$(customer\_id, employee\_id, type)$

$(employee\_id, branch\_name)$

$(customer\_id, branch\_name, employee\_id)$

- Observe that  $(customer\_id, employee\_id, type)$  contains a candidate key of the original schema, so no further relation schema needs to be added

- At end of for loop, detect and delete schemas, such as  $(employee\_id, branch\_name)$ , which are subsets of other schemas

- result will not depend on the order in which FDs are considered

- The resultant simplified 3NF schema is:

$(customer\_id, employee\_id, type)$

$(customer\_id, branch\_name, employee\_id)$

## → Decomposition to BCNF

→ conditions: For every FD  $\alpha \rightarrow \beta$  either:

①  $\alpha \rightarrow \beta$  is trivial ; or

②  $\alpha$  is a superkey

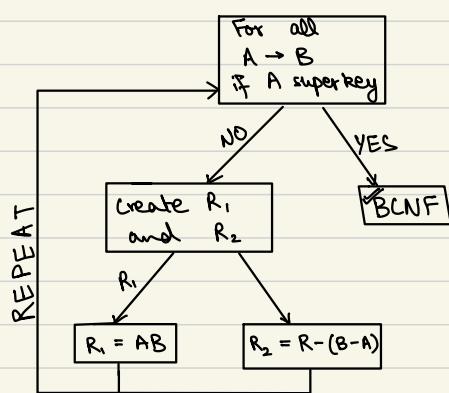
→ Testing for BCNF

→ no need to check on  $F^+$

→ to test if a decomposed relation is in BCNF, you need  $F^+$ , for testing on FDs (for all attributes in  $R_i$ ) that might not be present in  $F$  but can be logically implied.

→ testing for dependency preservation.  
↳ exponential or polynomial.

→ Algorithm:



- ① For all FD  $A \rightarrow B$  in  $F^+$ , check if  $A$  is superkey
- ② If  $A$  not superkey,  
create  $R_1 = A B$   
create  $R_2 = R - (B - A)$
- ③ Repeat for  $R_1$  and  $R_2$ .

## 3NF

① concentrates on Primary key.

② higher redundancy than BCNF.

③ FD  $X \rightarrow Y$  allowed if  $X$  is super key  
or  $Y$  is a part of some key.

## BCNF

concentrates on Candidate Key.

0% redundancy.

FD  $X \rightarrow Y$  allowed if  $X$  is super key.

## → Multivalued Dependency

- multi-valued FD notation:  $(\rightarrow)$  or  $(\rightarrow\rightarrow)$
- a table can be in BCNF and still have redundancy because of multi-valued dependencies.

## → Complementation

If  $X \rightarrow\rightarrow Y$ , then  $X \rightarrow (R - (X \cup Y))$

## Augmentation

If  $X \rightarrow\rightarrow Y$  and  $W \supseteq Z$ , then  $WX \rightarrow\rightarrow YZ$

## Transitivity

If  $X \rightarrow\rightarrow Y$  and  $Y \rightarrow\rightarrow Z$ , then  $X \rightarrow\rightarrow (Z - Y)$

## Replication

If  $X \rightarrow Y$ , then  $X \rightarrow\rightarrow Y$  but reverse untrue

## Coalescence

If  $X \rightarrow\rightarrow Y$ , and there is  $W$  s.t.  $W \cap Y = \emptyset$

$W \rightarrow Z$  and  $Y \supseteq Z$ , then  $X \rightarrow Z$

## → Decomposition to 4NF

→ If for all multi-valued dependencies in  $D^+$  of form  $\alpha \rightarrow\rightarrow \beta$  either:

- $\alpha \rightarrow\rightarrow \beta$  is trivial (i.e.,  $\beta \subseteq \alpha$  or  $\alpha \cup \beta = R$ )
- $\alpha$  is a superkey

→ If  $R$  is in 4NF then it is in BCNF.

## → Temporal Databases

→ time-varying data

→ have an associated time interval during which it is valid

→ Snapshot: value of data at a particular point in time

→ Timestamp Model

→ Two aspects of time:

Historical ① Valid time: time period which an attribute is valid in the real world.

Rollback ② Transaction time: time period for which the attribute is stored in the database.

→ Uni-temporal relation → either of the two aspects of time.

Bi-temporal " " → both aspects

→ Problem Solving on normalisation

Q. Let  $R(A, B, C, D)$  where all attributes have atomic values  
 $FD = \{B \rightarrow C, D \rightarrow A\}$ . Find the highest normal form  $R$  is in  
A.

Candidate keys = {BD}

All atomic values so 1NF ✓

$B \rightarrow C$  and  $B$  is proper subset of candidate key and  $C$  is non-prime  
2NF ✗

Q.  $R(A, B, C, D, E, F)$ .  $FD = \{BCD \rightarrow E, BCE \rightarrow F, DE \rightarrow A, DEA \rightarrow C, CA \rightarrow E\}$

A

Candidate keys = {BCD, BDE}

2NF ✗ because of BDE and  $DE \rightarrow A$

— X — X — X —

## WEEK 7

→ Characteristics of Application Programs

→ Applications are split into:

① **Frontend Layer** (presentation layer)

↳ interacts with the user

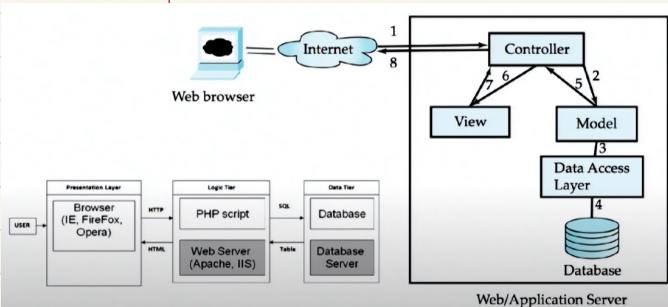
↳ example: choose item, checkout, etc.

→ MVC architecture

→ model: business logic

→ view: presentation of data

→ controller: receive actions, return view



② **Middle Layer** (application / business logic)

↳ functionality of application

↳ login authentication, search logic, etc.

→ high level view of data

→ hides details of data storage schema

③ **Backend Layer** (data access)

↳ manages persistent data

→ interfaces b/w business logic layer and underlying database.

- Business logic layer
  - ↳ abstraction of entities
  - ↳ gives you the object model
    - ↳ so you can use Java or C++
  - ↳ supports workflows

### ↳ Object - Relation Mapping

- looks at every entity as an object
- allows you to create an object of the relational model in the database so that you can use object-oriented programming

## → Tiered Architecture

- 1-tier : all components for an application on a single server
- 2-tier : client - server architecture
  - doesn't give separation b/w front-end and back-end
- 3-tier : Presentation , logic . and Data Access
  - most widely used to design a DBMS

→ URI

```

graph LR
 URI --> URL[URL → locator/address]
 URI --> URN[URN → name]

```

## → Sessions and cookies

- HTTP protocol is sessionless
  - ↳ i.e., once a server replies to a request , server closes the connection.
- Cookie : small piece of information
  - ↳ text containing identifying information.
  - ↳ can be stored permanently or limited time .

## → Java Server Pages

- similar concept as servlet
- compiled into Java + servlets
- much easier to write HTML embedding
- platform independent
- server-side
- can contain dynamic information unlike HTML

- Applications use API to interact with a database server
  - connect with server; send SQL commands; fetch tuples of result one-by-one

→ Frameworks

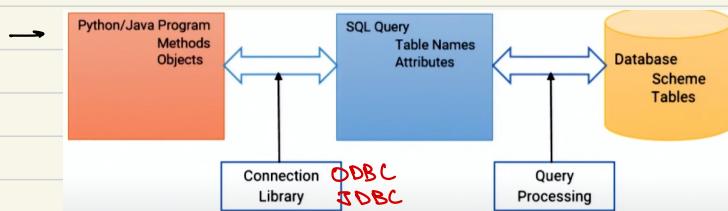
① Connectionist

- ODBC works with C, C++, Python
- JDBC works with Java
- SQL program is created, executed through a set of APIs.

② Embedding

- put SQL commands in your native language

→ Connectionist



→ ODBC

- Open Database Connectivity : standard API for accessing DBMS
- will work on all operating systems and database systems
- python example using pyodbc library →

```

import pyodbc
conn = pyodbc.connect('DSN=SQLS;UID=test01;PWD=test01')
cursor=conn.cursor()
cursor.execute("create table rvtest (col1 int, col2 float,
col3 varchar(10))")
cursor.execute("insert into rvtest values(1, 10.0,
'ABC')")
cursor.execute("select * from rvtest")

while True:
 row=cursor.fetchone()
 if not row:
 break
 print(row)

cursor.execute("delete from rvtest")
cursor.execute("insert into rvtest values (?, ?, ?)", 2,
20.0, 'XYZ')
cursor.execute("select * from rvtest")

while True:
 row=cursor.fetchone()
 if not row:
 break
 print(row)

```

→ JDBC

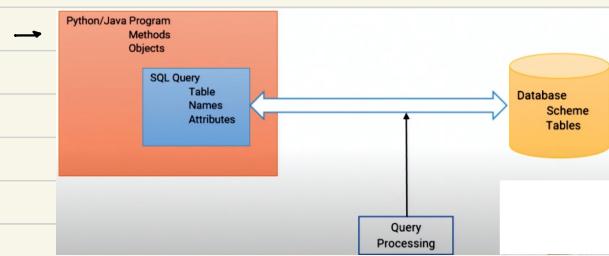
- Java Database Connectivity : API for the language Java , which defines how a client may access a database

→ Bridge Drivers

- bridges b/w JDBC and ODBC

- translates source function calls into target function calls

## → Embedding



→ Uses a tag and then puts the SQL statement inside it.  
**EXEC SQL <embedded SQL statement>**

→ EXEC SQL

*declare c cursor for output available in this cursor*

*select ID, name*

*from student*

*where tot\_cred > :credit\_amount* *variable from the host language*

**END\_EXEC**

**EXEC SQL open c ;**

**EXEC SQL fetch c into :si,:sn END\_EXEC**

**EXEC SQL close c ;**

→ PostgreSQL from Python

→ steps to access the database:

- ① Create connection
- ② Create cursor
- ③ Execute query
- ④ Commit / rollback
- ⑤ Close the cursor
- ⑥ Close the connection

→ example:

```

import psycopg2
def connectDb(dbname, username, pwd, address, portnum):
 conn = None
 try:
 # connect to the PostgreSQL database
 conn = psycopg2.connect(database = dbname,
 user = username,
 password = pwd,
 host = address,
 port = portnum)
 print("Database connected successfully")
 except (Exception, psycopg2.DatabaseError) as error:
 print(error)
 finally:
 conn.close()
connectDb('mydb', 'myuser', 'mypass', '127.0.0.1', "5432") # function call

import psycopg2
def createTable():
 conn = None
 try:
 # connect to the PostgreSQL database
 conn = psycopg2.connect(database = 'mydb',
 user = 'myuser',
 password = 'mypass',
 host = '127.0.0.1',
 port = '5432')
 cur = conn.cursor() # create a new cursor
 cur.execute('
CREATE TABLE EMPLOYEE
(
 emp_num INT PRIMARY KEY NOT NULL,
 emp_name VARCHAR(40) NOT NULL,
 department VARCHAR(40) NOT NULL
)') # execute the CREATE TABLE statement
 conn.commit()
 print("Table created successfully")
 cur.close()
 except (Exception, psycopg2.DatabaseError) as error:
 print(error)
 finally:
 if conn is not None:
 conn.close()
createTable() # function call

```

→ create a new table

## → Flask Web App Framework

- lightweight WSGI
- `app.route ("<URL>")`

↳ URL binding with the function

## → Rapid Application Development

- To speed up app development

→ function library to generate user-interface elements

→ drag-and-drop features in IDE to create UI elements (VS code)

→ automatically generate code for UI from a declarative spec

→ Phases:

- ① Business modeling
- ② Data modeling
- ③ Process modeling
- ④ Testing & turnover

→ Frameworks

### → Java Server Faces

→ a set of APIs for representing UI components

### → Ruby on Rails

→ easy creation of simple CRUD interfaces

→ Platforms and tools

G Suite, Google App Engine, Microsoft Azure, AWS Elastic

## → Application Performance

→ caching techniques to reduce cost of serving pages.

→ can be at server site or at the client's network.

## → Application Security

→ SQL injection use prepared statements with user inputs as parameters

→ Never store passwords in clear texts, encrypt it very well

→ restrict access to database by source IP address

→ Authentication

→ password plus OTP

→ face recognition, fingerprint, etc.

→ Audit trail at the database and the application level

→ are used after-the-fact

## → Types of Mobile Apps:

① Native : native language of a platform.

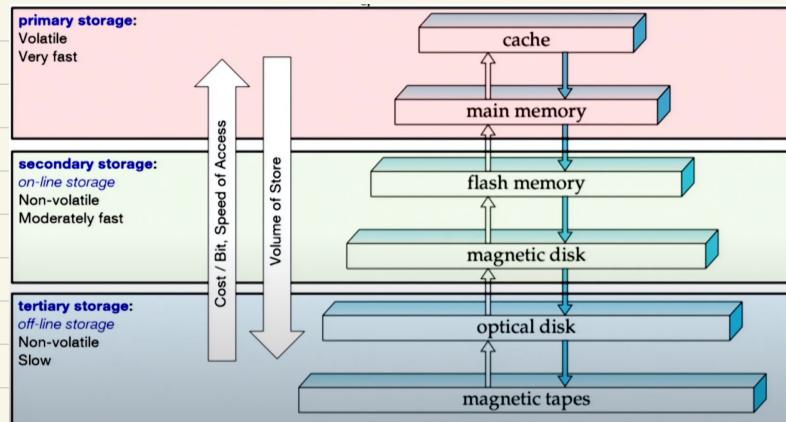
② Web : runs completely inside of a Web browser

③ Hybrid : combines attributes of both.



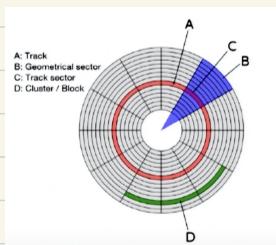
# WEEK 8

- **Algorithm** : finite sequence of well-defined, computer-implementable instructions.
  - must terminate
- **Program** : collection of instructions
  - may or may not terminate
- **Physical Storage Media**
  - speed
  - cost per unit of data
  - reliability
  - **Volatile storage** : loses contents when power is switched off
    - **cache** : very fast and most costly storage
    - **main memory** : fast access ; very small
  - **Flash memory** : data survives power failure
    - reads as fast as main memory
  - **Magnetic Disk** : primary medium for long-term storage
    - direct-access : possible to read data on disk in any order, unlike magnetic tape
    - survives power failure
  - **Optical Storage** : write once and read many times (archival data)
  - **Tape Storage** : non-volatile ; for backup and archival
    - sequential access
    - very high capacity



## → Magnetic Disk

### → Mechanism



- surface divided into circular tracks
- each track divided into sectors

→ **Disk controller** : interfaces b/w computer system and disk drive hardware

↳ computes and attaches **checksums** to each sector to verify that correct read back

→ Access time : **Seek time** and **Rotational latency**

↳ **Seek time** : time to reposition the arm over the correct track

↳ **Rotational latency** : to get to that particular sector

→ Data-transfer rate

↳ typically 25 to 100 MB per second

→ Mean time to failure (MTTF)

↳ typically 3-5 years

## → Magnetic Tapes

→ large volumes of data but lesser than disks

→ very cheap

→ limited to sequential access

## → Cloud Storage

| Parameters       | Cloud Storage                                                                                                                                                                                                     | Traditional Storage                                                                                                                |
|------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------|
| Cost             | Cloud storage is cheaper per GB than using external drives.                                                                                                                                                       | The hardware and infrastructure costs are high and adding on more space and upgrading only adds extra costs.                       |
| Reliability      | Cloud storage is highly reliable as it takes less time to get under functioning                                                                                                                                   | Traditional storage requires high initial effort and is less reliable.                                                             |
| File Sharing     | Cloud storage supports file sharing dynamically as it can be shared anywhere with network access                                                                                                                  | Traditional storage requires physical drives to share data and a network is to be established between both                         |
| Accessibility    | Cloud storage gives you access to your files from anywhere                                                                                                                                                        | Restricted to local access                                                                                                         |
| Backup/ Recovery | Very safe from on site disaster. In case of a hard drive failure or other hardware malfunction, you can access your files on the cloud, which acts as a backup solution for your local storage on physical drives | Data that is stored locally is much more susceptible to unexpected events and local storage and local backups could be easily lost |

## → Future of Storage

→ DNA digital storage

→ Quantum memory

## → File Organisation

→ Database = collection of files

File = sequence of records

Record = sequence of fields

→ database file is partitioned into fixed-length storage units called **blocks**

→ Deletion of a record i : Alternatives

① move records  $i+1, \dots, n$  to  $i, \dots, n-1$

② move record  $n$  to  $i$

③ do not move records but link all free records on a **free list**.

→ Variable-length records

→ storage of multiple record types in a file

→ record types that allow variable length for fields

→ **slotted page** header contains:

→ number of record entries

→ end of free space in a block

→ location and size of each record

## → Organisation of records in files

→ **Heap** : record placed anywhere in the file where there is space

**Sequential** : in sequential order based on value of search key of each record

**Hashing** : hash function computed on some attribute of each record

## → Sequential file organisation

↳ deletion : use pointer chains

insertion : locate position ; if no free space insert into **overflow block**

↳ need to reorganise to restore the order

## → Multitable clustering

↳ Store several relations in one file using a **multitable clustering** file organization

| dept_name  | building | budget |
|------------|----------|--------|
| Comp. Sci. | Taylor   | 100000 |
| Physics    | Watson   | 70000  |

| dept_name  | building | budget |
|------------|----------|--------|
| Comp. Sci. | Taylor   | 100000 |
| Physics    | Watson   | 70000  |

| ID    | name       | dept_name  | salary |
|-------|------------|------------|--------|
| 10101 | Srinivasan | Comp. Sci. | 65000  |
| 33456 | Gold       | Physics    | 87000  |
| 45565 | Katz       | Comp. Sci. | 75000  |
| 83821 | Brandt     | Comp. Sci. | 92000  |

| dept_name  | building   | budget |
|------------|------------|--------|
| Comp. Sci. | Taylor     | 100000 |
| 45564      | Katz       | 75000  |
| 10101      | Srinivasan | 65000  |
| 83821      | Brandt     | 92000  |
| Physics    | Watson     | 70000  |
| 33456      | Gold       | 87000  |

group records of  
instructors from the same  
department

→ good for queries involving department & instructor

## → Data Dictionary Storage

→ data dictionary stores **metadata**

- ↳ information about relations
- ↳ statistical and descriptive data
- ↳ physical file organisation information

## → Storage Access

→ minimise number of block transfers b/w disk and memory

→ **buffer**: portion of main memory available to store copies of disk blocks

→ **buffer manager**: subsystem responsible for allocating buffer space

↳ programs call on buffer manager when they need a block from disk

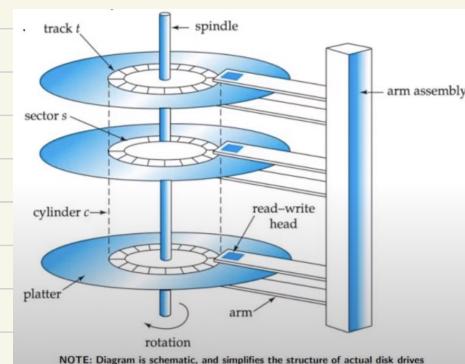
↳ **LRU strategy**: Least Recently Used → use past patterns of block references as a predictor of future references

↳ **Pinned block**: memory block that is not allowed to be written back to disk

↳ **loss-immediate strategy**: free the space occupied by a block as soon as the last tuple has been processed.

↳ **MRU strategy**: Most Recently Used → pin the block currently being used; unpin the block once processed.

## → Tutorial



→ sector size : 512 byte

→ number of cylinders = number of tracks per surface

— X — X — X —

# WEEK 9

## → Indexing

- use a separate table for 'indexed' attribute.  
The index table is sorted on the attribute and also has an additional attribute 'pointer' that points to the tuples in original data.
- kinds of indices
  - ↳ ordered indices
  - ↳ hash indices
- index evaluation : access types supported ; access time ; insertion time ; deletion time ; space overhead

## → Ordered Indices

- entries stored in sorted order on the search key (indexed attribute) value.
- primary index (a.k.a. clustering index)
  - ↳ in a sequentially ordered file , the index whose search key specifies the sequential order of the file
  - ↳ usually (but not always) primary key .

## → secondary index (a.k.a. non-clustering index)

- ↳ index whose search key specifies a diff. order from the sequential order of the file .

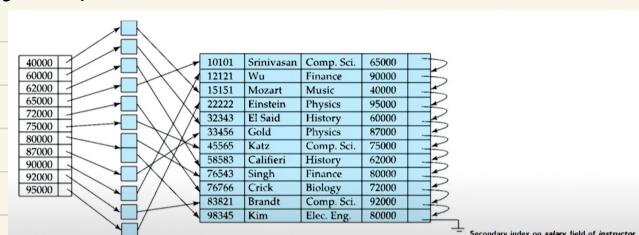
## → dense index

- ↳ index record appears for every search-key value in the file .
- ↳ might become very large

## → sparse index

- ↳ index records for only some search-key values
- ↳ to find record with search-key value  $k$  , find index record with largest search-key value  $< k$
- ↳ first record of every block as index entry .

- if secondary index has repeated values , maintain only distinct
  - ↳ index record points to a bucket that contains pointers
  - ↳ has to be dense



## → multilevel index

↳ if primary index does not fit in memory

↳ solution: treat primary index file as sequential file and construct a sparse index on it.

outer index: sparse index on primary index

inner index: primary index file

## → Deletion in index

↳ dense: delete similar search-key in index table

↳ sparse:

→ if entry of the search-key exists in the index table, delete it and replace with the next search-key value in the file (in search-key order)

## → Insertion in index

↳ dense: if search-key value doesn't appear in index, insert it

↳ sparse:

→ if a new block is created, the first search-key in the new block is inserted into the index

## → Balanced Binary Search Trees

→ BST is balanced if height  $\sim \log(n)$

→ Balancing strategies:

→ AVL tree

→ Skip List

→ Randomised BST

→ Splay

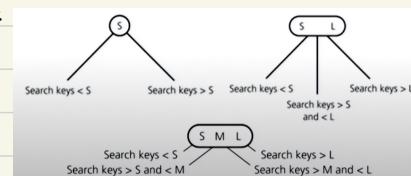
→ These strategies work well for in-memory operations, small volume of data but have complex operations and do not scale for external data

## → 2-3-4 Tree

→ guarantee that all leaves are at same depth (bottom level)

↳ height of all leaf nodes are same;  $h \sim O(\log n)$

→ use three kinds of nodes:



→ insertion: search to find expected location:

→ if 2 node: change to 3 node and insert

→ if 3 node: change to 4 node and insert

→ if 4 node: split the node by moving the middle item into parent node, and insert

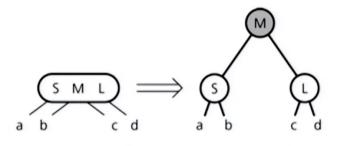
→ node splitting if :

↳ 4-node is the root ; or

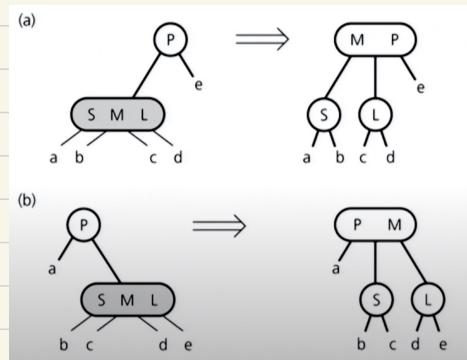
↳ 4-node has a 2-node parent ; or

↳ 4-node has a 3-node parent

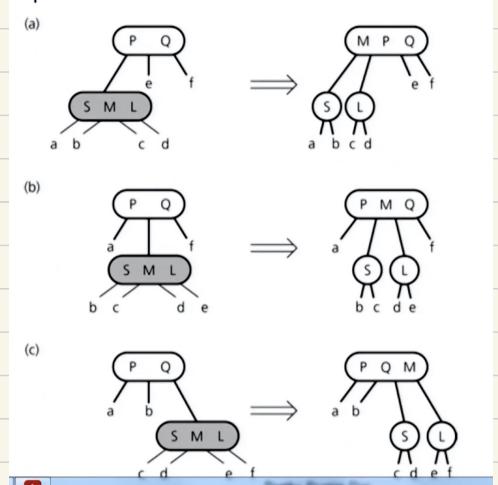
→ splitting at root



→ splitting with 2-node parent



→ splitting with 3-node parent



→ height only changes when the split happens at the root  
but the height for all leaf nodes is still equal.

→ strategies

early → split a 4-node as soon as you come across it in traversal ; ensures that the path doesn't have multiple 4-nodes.

late → split a 4-node only when you need to insert an item in it ; could lead to  $O(n)$  splits

↳ both strategies have the same complexity.

## → B<sup>+</sup> Tree Index Files

→ B<sup>+</sup> Tree : balanced binary search tree

↳ follows multi-level index like 2-3-4 tree

↳ leaf nodes denote actual data pointers

↳ leaf nodes are linked using a link list

↳ **internal nodes**

→ at least  $n/2$  child pointers except root

→ at most  $n$  pointers

↳ **leaf nodes**

→ at least  $n/2$  record pointers and  $n/2$  key values

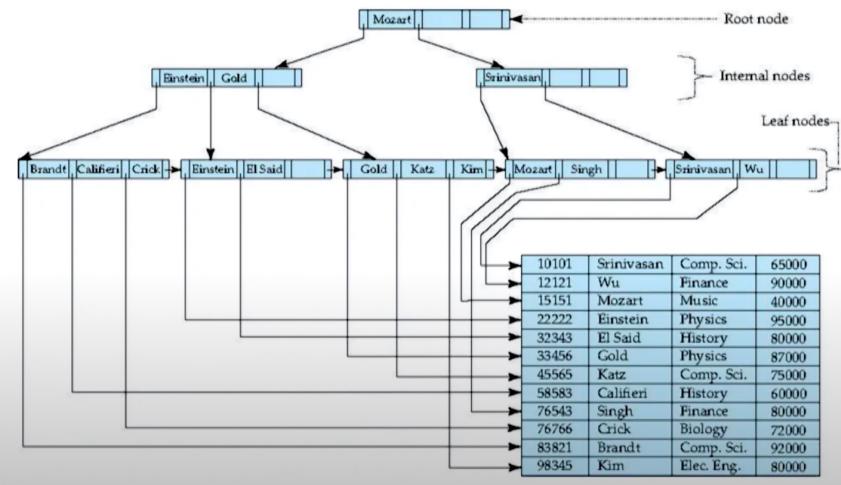
→ at most  $n$  " " " " " "

↳ **insertion**

→ if the target location is full, split the node  
following the process as in 2-3-4 tree.

↳ re-organises itself, as opposed to Index Sequential file  
↳ disadvantage : extra insertion and deletion overhead

↳ **example :**

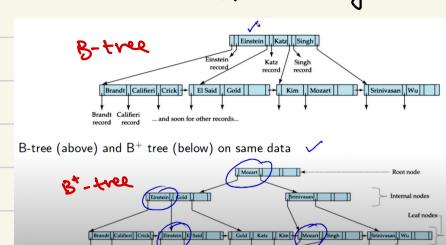


↳ non-leaf levels form a hierarchy of sparse indices  
↳ since every node must have  $n/2$  values, to search for a key value, it takes  $O(\log_{n/2} K)$ , where  $K = \text{total } \# \text{ of keys}$ .

↳ can also be used for database files.

## → B-Tree

- similar to B<sup>+</sup> tree but allows search-key values to appear only once.
- uses less tree nodes
- sometimes possible to find before reaching leaf nodes
- non-leaf nodes are larger ; height increases
- insertion and deletion becomes more complicated.



## → Static Hashing

### → Hash Function

↳ maps data of arbitrary size (from domain D) to fixed-size values (say, integers from 0 to  $N > 0$ ).

$$h : D \rightarrow [0, \dots, N]$$

→ bucket : unit of storage containing one or more records

→ records with diff. search-key values may be mapped to the same bucket so the entire bucket has to be searched sequentially

→ ideal hash function is uniform : each bucket is assigned some number of search-key values

→ ideal hash function is random

→ bucket overflows reasons :

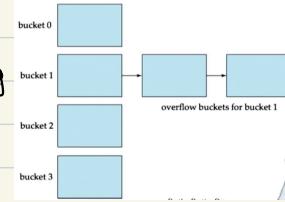
→ insufficient buckets

→ skew in distribution of records

handling:

→ overflow buckets

↳ closed hashing



→ problematic because of fixed number of buckets.

## → Dynamic Hashing

→ extendable hashing ; hash function generates values over a large range → typically b-bit integers with  $b = 32$

→ use only a prefix of hash function to index into a table of bucket addresses.

↳ let length of prefix be i bits ,  $0 \leq i \leq 32$

→ bucket address table size =  $(2)^i$

→ value of i grows and shrinks with the database

↳ example with  $b=3$

① 0 0 0      ⑤ 1 0 0

② 0 0 1      ⑥ 1 0 1

③ 0 1 0      ⑦ 1 1 0

④ 0 1 1      ⑧ 1 1 1

with  $i=1$  , bucket 0 = {①, ②, ③, ④}

bucket 1 = {⑤, ⑥, ⑦, ⑧}

with  $i=2$  , bucket 0 = {①, ②}      2 = {⑤, ⑥}

bucket 1 = {③, ④}      3 = {⑦, ⑧}

### → insertion :

- search : get to index , get the hash , take the prefix , go to the bucket address table , then go to the bucket.
- if target bucket is full : options
  - ① allocate a new bucket
  - ② if you can't create anymore buckets , create an overflow bucket OR use more bits

### → Bitmap Index

- for attributes that have a limited number of possible values , like gender
- has as many bits as records
- in a bitmap for value  $v$  , the bit of a record is 1 if the record has the value  $v$  for the attribute , and is 0 otherwise

### → Index Definition in SQL

- create index <index-name> on <relation-name> (<attribute-list>)
- Tablespace → can specify which tablespace to go to

#### → STORAGE → how to store a database object

- ↳ INITIAL size of the 1<sup>st</sup> extent of object
- ↳ NEXT " " " 2<sup>nd</sup> " " "
- ↳ PCTINCREASE how much %age later extents grow
- ↳ PCTFREE % of each data block be free for updates

#### → Multiple keys

- ↳ multiple indices ; composite index
- ↳ composite search key (order is important)
  - lexicographic ordering :  $(a_1, a_2) < (b_1, b_2)$  if either
    - $a_1 < b_1$  , or
    - $a_1 = b_1$  and  $a_2 < b_2$

#### → To create an index

- ↳ you must own , or have the INDEX object privilege

## → Guidelines For Indexing

- collect statistics about data to learn the patterns and adjust indexes
- **Ground Rules**:

- ① Indexes lead to Access-Update tradeoff
- ② Index the correct tables
- ③ Index the correct columns
- ④ Limit the number of indexes for each table
- ⑤ Choose the order of the columns in composite indexes
- ⑥ Gather statistics about usage
- ⑦ Drop indexes no longer required

— X — X — X —

# WEEK 10

## → Transactions

- unit of program execution that accesses and possibly updates various data items

### → ACID properties

#### → **Atomicity** (All or nothing transactions)

↳ ensure updates of partial transaction are not in database

#### → **Consistency** (Guarantees committed transaction state)

↳ explicitly specified integrity constraints (primary / foreign keys)

↳ implicit integrity constraints

#### → **Isolation** (Transactions are independent)

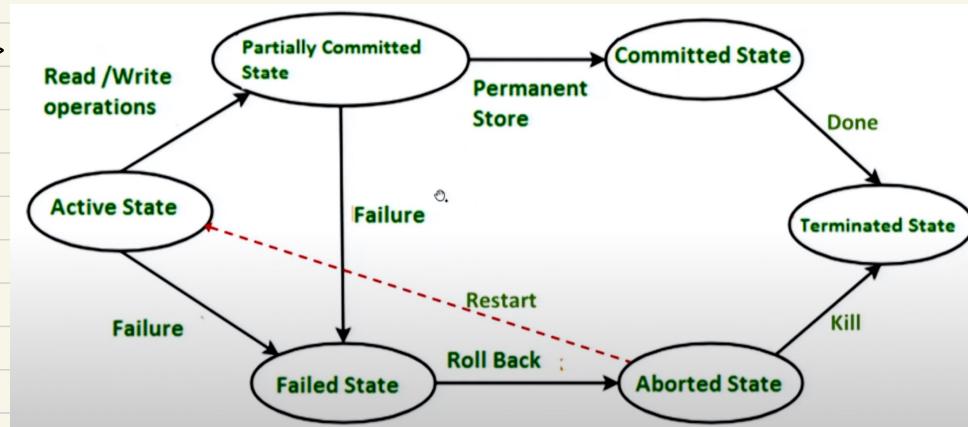
↳ if multiple transactions, no inconsistent updates are made together

↳ run transactions serially

#### → **Durability** (committed never lost)

↳ once a transaction is committed, it will remain committed

## → States ⇒



## → Concurrent Executions

- multiple transactions allowed to run concurrently
  - ↳ increased disk utilization
  - ↳ reduced average response time
- Schedule: sequence of instructions
  - ↳ must consist all transactions
  - ↳ must preserve order of instructions within transaction
- successful (failed) transaction will have commit (abort) at the end.
- in a schedule, the database might be inconsistent state at one of the commits but that is acceptable as long as the database is in a consistent state after the last commit.
- all schedules may not satisfy ACID

## → Serializability

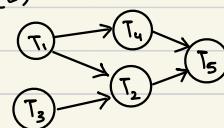
- assumption: each transaction preserves db consistency
- schedule is serializable if equivalent to a serial schedule
  - ↳ conflict serializability
  - ↳ view serializability
- conflicting instructions
  - ↳ Two instructions  $I_1$  and  $I_2$  accessing the same item
    - if  $I_1 = \text{'write'}$  or  $I_2 = \text{'write'}$ :  
 $I_1$  and  $I_2$  are conflicting
    - if they're consecutive no conflict

## → Conflict Serializability

- Transform schedule  $S$  into  $S'$  by swapping non-conflicting instruction.  $S$  and  $S'$  are conflict equivalent
- Schedule  $S$  is conflict serializable if it is conflict equivalent to a serial schedule.
- If a schedule is serializable, it doesn't necessarily mean that it is conflict serializable.
  - ↳ ex.  $w_1(A), w_2(A), w_2(B), w_1(B), w_3(B)$
  - serial:  $w_1(A), w_1(B), w_2(A), w_2(B), w_3(B)$

## → Precedence Graph

- ↳ nodes: transactions
  - ↳ if two transactions conflict then they're connected
- ↳ a schedule is conflict serializable iff the precedence graph is acyclic.
  - ↳ serializability can be obtained by topological sorting
- ↳ example:  $w_1(A), r_2(A), w_1(B), w_3(C), r_2(C), r_4(B), w_2(D), w_4(E), r_5(D), w_5(E)$



$T_1, T_3, T_4, T_2, T_5$

## → Recovery

- if transaction  $T_j$  reads an item previously written by  $T_i$ , then the commit operation of  $T_i$  must appear before commit of  $T_j$ 
  - ↳ then this schedule is recoverable in case of a system crash

## → cascading rollback

↳ example →

| $T_{10}$                          | $T_{11}$              | $T_{12}$ |
|-----------------------------------|-----------------------|----------|
| read (A)<br>read (B)<br>write (A) |                       |          |
|                                   | read (A)<br>write (A) |          |
|                                   |                       | read (A) |

If  $T_{10}$  fails, ↪ abort

$T_{11}$  and  $T_{12}$  must also be rolled back

— schedule is recoverable because no transaction commits

## → cascadeless schedule

- ↳ For each  $T_i$  and  $T_j$ ,  $T_j$  reads an item written by  $T_i$ , the commit of  $T_i$  appears before read operation of  $T_j$
- ↳ every cascadeless is also recoverable

Example: Irrecoverable Schedule

| T1            | T1's Buffer | T2           | T2's Buffer | Database |
|---------------|-------------|--------------|-------------|----------|
| R(A);         | A = 5000    |              |             | A = 5000 |
| A = A - 1000; | A = 4000    |              |             | A = 5000 |
| W(A);         | A = 4000    |              |             | A = 4000 |
|               |             | R(A);        | A = 4000    | A = 4000 |
|               |             | A = A + 500; | A = 4500    | A = 4000 |
|               |             | W(A);        | A = 4500    | A = 4500 |
|               |             | Commit;      |             |          |
| Failure Point |             |              |             |          |
| Commit;       |             |              |             |          |

Example: Recoverable Schedule with Cascading Rollback

| T1            | T1's Buffer | T2           | T2's Buffer | Database |
|---------------|-------------|--------------|-------------|----------|
| R(A);         | A = 5000    |              |             | A = 5000 |
| A = A - 1000; | A = 4000    |              |             | A = 5000 |
| W(A);         | A = 4000    |              |             | A = 4000 |
|               |             | R(A);        | A = 4000    | A = 4000 |
|               |             | A = A + 500; | A = 4500    | A = 4000 |
|               |             | W(A);        | A = 4500    | A = 4500 |
| Failure Point |             |              |             |          |
| Commit;       |             |              |             |          |
|               |             | Commit;      |             |          |

Example: Recoverable Schedule without Cascading Rollback

| T1            | T1's Buffer | T2           | T2's Buffer | Database |
|---------------|-------------|--------------|-------------|----------|
| R(A);         | A = 5000    |              |             | A = 5000 |
| A = A - 1000; | A = 4000    |              |             | A = 5000 |
| W(A);         | A = 4000    |              |             | A = 4000 |
| Commit;       |             | R(A);        | A = 4000    | A = 4000 |
|               |             | A = A + 500; | A = 4500    | A = 4000 |
|               |             | W(A);        | A = 4500    | A = 4500 |
|               |             | Commit;      |             |          |

## → Transaction Definition is SQL

→ transaction end by: Commit work or by Rollback work

→ TCL Transaction Control Language

COMMIT → save changes

ROLLBACK → roll back change

SAVEPOINT → points within groups of transactions to rollback

SET TRANSACTION → names a transaction

↳ are only used with DML commands: INSERT, UPDATE, DELETE

## → View Serializability

→ weaker than conflict serializability

→ S and S' schedules are view equivalent if:

### ① Initial Read

↳ If  $T_i$  in S reads initial value of Q, then in S' also  $T_i$  must read initial value of Q.

### ② Write - Read Pair

↳ If  $T_i$  in S executes read(Q) which reads the value written by  $T_j$ , then in S' this order of write-read must be maintained.

### ③ Final Write

↳ The transaction that performs the final write(Q) operation in S must also be the last one in S'.

→ view serializable example:

(but not conflict serializable)

| $T_{27}$ | $T_{28}$ | $T_{29}$ |
|----------|----------|----------|
| read(Q)  |          |          |
| write(Q) | write(Q) | write(Q) |

$T_{28}$  and  $T_{29}$  write(Q)  
are blind writes

## → Concurrency Control

→ develop concurrency control protocols that will assure serializability

→ one way to ensure isolation: require that data items be accessed in a mutually exclusive manner.

→ allow a transaction to access a data item only if it is holding a lock on that data item.

## → Lock-Based Protocol

→ data items locked in two modes:

### ① exclusive (X) mode

↳ both read and write

↳ using lock-X instruction

### ② shared (S) mode

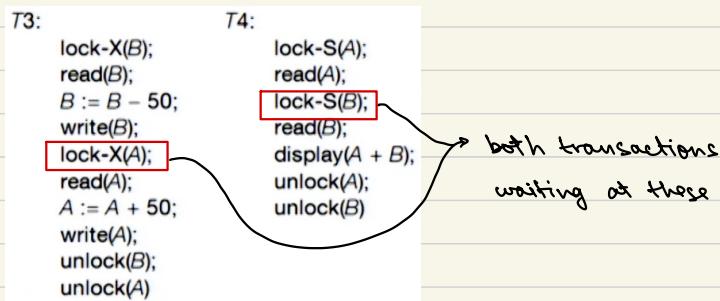
↳ only read

↳ using lock-S instruction

→ unlock data item Q by unlock-S instruction

→ Deadlock when neither transaction can ever proceed because of a 'loop' in locking of data items.

↳ example:



both transactions will be waiting at these points!

↳ the system must roll back one of the transactions.

→ Two-Phase Locking Protocol

↳ Phase 1: Growing Phase

↳ transaction obtains but not release locks

Phase 2: Shrinking Phase

↳ release locks but not obtain

↳ assures serializability ; in the order of lock points

→ Lock Conversions

↳ Phase 1

↳ can lock-S on item

" lock-X on item

" convert lock-S to lock-X

Phase 2

↳ can release lock-S

" " lock-X

" convert lock-X to lock-S

↳ doesn't ensure freedom from deadlock

→ Starvation (Livelock)

↳ when a transaction is trying to get an X-lock while a sequence of other transactions request and are granted an S-lock on the item.

→ Strict Two-phase protocol

↳ transaction must hold all its X-locks till it commits / aborts.

↳ Rigorous → all locks held until commit / abort

## → Deadlock Handling

→ strategies:

- each transaction locks all its data items before execution
- partial ordering

→ Transaction timestamp : created by DBMS to identify relative starting time

① Wait-die scheme : non-preemptive

↳ older transactions wait for younger ones

↳ ex.  $T_5$  requests data held by  $T_{10}$  will wait

$T_5$  requests " " " " " be killed

② Wound-wait scheme : preemptive

↳ older transactions forces roll back of younger ones

↳ ex.  $T_5$  requests data held by  $T_{10}$ ,  $T_{10}$  will be wounded

$T_5$  requests " " " " ,  $T_5$  will wait

→ Deadlock detection

↳ graph with transactions as nodes.

If  $T_i$  is waiting for  $T_j$ , then there is an edge from  $T_i$  to  $T_j$

↳ if graph has cycle, it has a deadlock.

→ Deadlock Recovery

↳ most effective to rollback to the lockpoint that has caused the deadlock.

↳ include number of rollbacks in the cost factor to avoid starvation.

## → Timestamp Based Protocols

→ maintains for each data Q:

→ W-timestamp (Q) : largest successful write (Q)

→ R-timestamp (Q) : " " " read (Q)

→  $T_i$  issues a read (Q):

if  $TS(T_i) \leq W\text{-timestamp}(Q)$ : read rejected and  $T_i$  rolled back  
if  $TS(T_i) \geq W\text{-timestamp}(Q)$ : read accepted and R-timestamp(Q) updated

→  $T_i$  issues a write (Q):

if  $TS(T_i) < R\text{-timestamp}(Q)$ : write rejected ;  $T_i$  rolled back  
if  $TS(T_i) < W\text{-timestamp}(Q)$ : " " " " "  
else allow write operation and update W-timestamp (Q)

→ no cycles in precedence graph

→ ensures deadlock-free

→ may not be cascadeless or recoverable



# WEEK 11

## → Backup and Recovery

→ Backup: representative copy of data

↳ Physical Backup

↳ copy of physical database files

↳ e.g. data, control files, log files

↳ Logical Backup

↳ copy of logical data extracted from db

↳ e.g. tables, procedures, views, functions

→ uses of backup

↳ disaster recovery

↳ client side changes

↳ auditing

↳ downtime

→ Types of backup data

↳ Business data: personal info of clients, employees

↳ System data: environment / configuration of system

↳ Media: photos, videos, graphics

## → Backup Strategies

→ Full Backup

↳ backs up everything

↳ must be done at least once before any other type of backup

↳ frequency depends on application

→ Incremental Backup

↳ backup only files that have changed since the last backup

↳ cannot be done without full backups

→ Differential Backup

↳ backup only files that have changed since the last full backup

## → Hot Backup

- preferable wherever possible
  - refer to keeping a database up and running while the backup is performed concurrently
  - database is always available
  - detecting error can be difficult
- mainly used for Transaction Log Backup

## → Failure Classification

- Transactions have ACID properties
- Transaction failure
  - logical errors
  - system errors
- System crash
- Disk crash

## → Recovery Algorithms

- ↳ actions taken during transaction processing to ensure enough information exists to recover from failures
- ↳ actions taken after a failure to recover db contents to state that ensures atomicity, consistency, durability

## → Storage Structure

- Volatile : does not survive system crash
  - ↳ main memory , cache memory
- non volatile : survives system crash
  - ↳ disk , tape , flash memory
- Stable storage : mythical form that survives all failures
  - ↳ multiple copies on distinct non-volatile media

## → Stable Storage implementation

- ↳ assuming two copies of each block

↳ the output is complete only after 2<sup>nd</sup> write completes

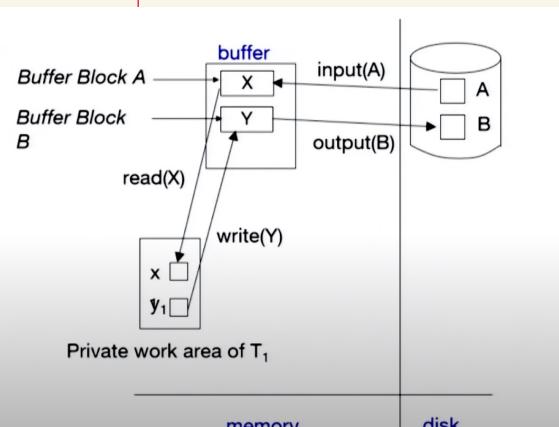
## → Failure during data transfer.

- ↳ find inconsistent blocks

↳ record in-progress disk writes and use this information during recovery to find inconsistent blocks

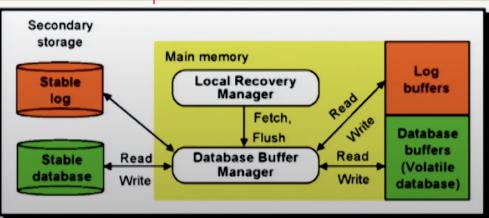
↳ if either copy has an error , overwrite it by the other

↳ if both have no error , overwrite 2<sup>nd</sup> block by the 1<sup>st</sup> block



## → Log-Based Recovery

- Log is created for every transaction
- When a transaction  $T_i$  starts, it registers itself by writing a record  $\langle T_i \text{ start} \rangle$  to the log.
- ↳ when  $T_i$  executes  $\text{write}(x)$ , log record  $\langle T_i, x, V_1, V_2 \rangle$  is written, where  $V_1$  = old value and  $V_2$  = new value
- ↳ Finally  $T_i$  finishes last statement.  $\langle T_i \text{ commit} \rangle$  recorded



## → Database Modification Schemes

- ↳ Immediate Modification allows updates of uncommitted transaction to be made to the buffer, or the disk itself
- ↳ Deferred Modification performs updates to buffer/disk only at the time of transaction commit.

→ Transaction is said to have committed when its commit log record is output to stable storage.

- Undo of a log record  $\langle T_i, x, V_1, V_2 \rangle$  writes the old value  $V_1$  to  $x$
- ↳ Redo - - - - " " " " " " new "  $V_2$  to  $x$
- Undo is used for rollback during normal operation
- Undo and Redo are used during recovery from failure

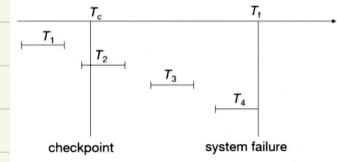
## → Checkpoints

↳ to streamline recovery

↳ All updates are stopped while doing checkpointing

- ① Output all log records from main memory to storage
- ② Output all modified buffer blocks to disk
- ③ Write log record  $\langle \text{checkpoint } L \rangle$  onto storage where  $L$  is a list of all transactions active at the time of checkpoint

↳ example:



$T_1$  can be ignored.

$T_2$  and  $T_3$  redone.

$T_4$  undone

commits before checkpoint: ignored

commits after checkpoint: redone

uncommitted transactions: undone

## → Transactional Logging

→ Hot backup

↳ take a backup while doing transactions

→ Recover: retrieve from the backup media .db files, and transactions log

Restore: re-apply db consistency based on the transaction log

## → Recovery Algorithm

→ assume if a transaction  $T_i$  has modified an item, no other transaction can modify the same item until  $T_i$  has committed or aborted.

→ Logging (normal operation):

<  $T_i$  start >

<  $T_i, X, V_1, V_2$  >

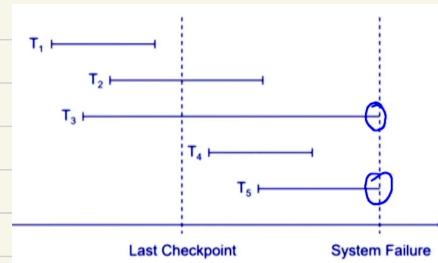
<  $T_i$  commit >

→ Transaction rollback (normal operation):

↳ scan the log backwards and undo each log record

→ For <  $T_i, X, V_1, V_2$  >, undo  $X$  to  $V_1$ , and write a log record <  $T_i, X, V_1$  > Compensation Log Record

→ example:



- $T_1$  ignored
- $T_2$  and  $T_4$  redone
- $T_3$  and  $T_5$  undone and restarted

→ Redo phase:

① Find last < checkpoint L > record, and set undo-list L

② Scan forward from < checkpoint L > record

→ whenever record <  $T_i, X, V_1, V_2$  > found, redo it

→ " " <  $T_i$  start > found, add  $T_i$  to undo-list

→ " " <  $T_i$  commit > or <  $T_i$  abort > found, remove  $T_i$  from undo-list

→ At the end of this phase, undo-list will have all the transactions that need to be undone.

→ Then do the **undo phase** (reverse each action).

## → Recovery with Early Lock Release

- any index used in processing a transaction, such as a B<sup>+</sup>-tree, can be treated as normal data.
- B<sup>+</sup>-tree concurrency algorithm allows locks to be released early
  - ↳ because of early release it is possible:
    - T<sub>1</sub> updates a value in B<sup>+</sup> tree by inserting entry (V<sub>1</sub>, R<sub>1</sub>)
    - T<sub>2</sub> inserts (V<sub>2</sub>, R<sub>2</sub>) in the same node, moving the entry (V<sub>1</sub>, R<sub>1</sub>) before completion of T<sub>1</sub>.
  - ↳ we cannot undo T<sub>1</sub> since that would also undo the insert done by T<sub>2</sub>
  - ↳ the only way to undo the effect of insertion (V<sub>1</sub>, R<sub>1</sub>) is to execute a deletion.

→ Logical undo

## → Operation Logging : Process

- ↳ Operation starts → log < T<sub>i</sub>, O<sub>j</sub>, operation-begin >
- ↳ while executing operation, log records in the normal fashion
- ↳ Operation completes → < T<sub>i</sub>, O<sub>j</sub>, operation-end, U>
  - contains information needed to perform a logical undo

### ↳ example:

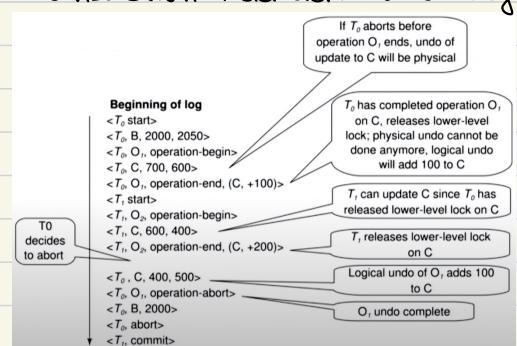
- Insert of (key, record-id) pair (K5, RID7) into index I9

```

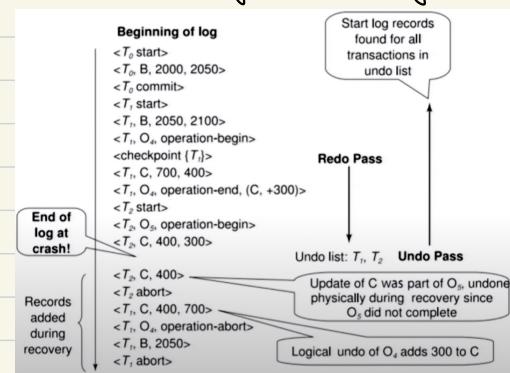
<T1, O1, operation-begin>
...
<T1, X, 10, K5>
<T1, Y, 45, RID7>
}
<T1, O1, operation-end, (delete I9, K5, RID7)>

```

## → example of Transaction Rollback with Logical Undo



## → example of Failure Recovery with Logical Undo



## → Plan for Backup and Recovery

- Data Importance
- Frequency of change
- Speed of backup or recovery
- Equipment
- Employees
- Storing

## → RAID: Redundant Array of Independent Disks

- disk organisation techniques that manage a large number of disks providing a view of a single disk with:
  - ↳ high capacity and high speed
  - ↳ high reliability

### → Mirroring

- ↳ duplicate every disk
- ↳ every write done on both disks
- ↳ read from either
- ↳ data loss would occur only if one disk fails and its mirror disk also fails before the system is repaired

### → Bit-level Striping

- ↳ splits bits of each byte across multiple disks
- ↳ seek / access time worse than for a single disk

### → Byte-level Striping

- ↳ each file split into parts one byte in size

### → Block-level Striping

- ↳ with  $n$  disks, block  $i$  of a file goes to disk  $(i \bmod n) + 1$
- ↳ requests for diff. block can run in parallel if they reside on different disks

### → Bit-Interleaved Parity

- ↳ single parity bit enough for error detection and correction

### → Block-Interleaved Parity

- ↳ uses block-level striping and keeps a parity block on a separate disk for corresponding blocks from  $n$  other disks

### → Standard RAID levels

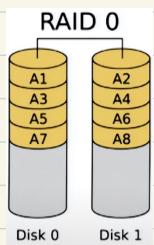
- ↳ RAID 0: striping ; RAID 1: mirroring ; RAID 5: distributed parity;

RAID 6: dual parity

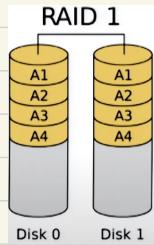
- ↳ can combine (or nest) levels: RAID 01: mirroring stripe sets

RAID 10: striping of mirrors

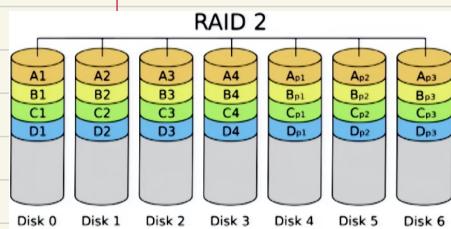
- RAID only a building block of a data loss prevention and recovery scheme but cannot replace a backup plan



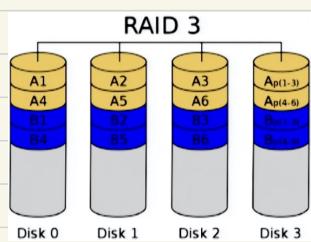
→ RAID 0 : striping  
 ↳ no redundant data  
 ↳ least costly  
 ↳ reliability is poor  
 ↳ best write performance



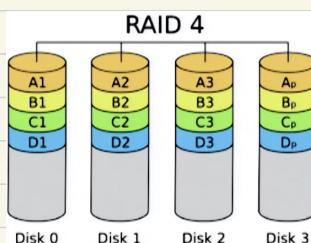
→ RAID 1 : Mirroring  
 ↳ two identical copies  
 ↳ most expensive  
 ↳ excellent fault tolerance  
 ↳ space utilisation 50 percent



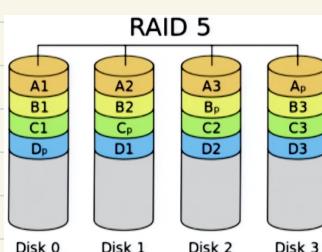
→ RAID 2 : Parity  
 ↳ striping unit is single bit  
 ↳ Hamming code for parity  
 ↳ For a 4-bit data, 3 bits are added



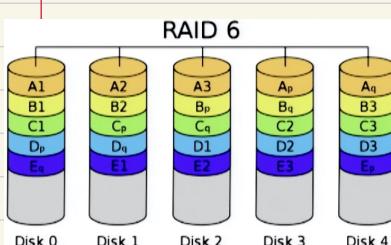
→ RAID 3 : Byte Striping + Parity  
 ↳ single check disk with parity information  
 ↳ cannot service multiple requests simultaneously



→ RAID 4 : Block Striping + Parity  
 ↳ striping unit of a disk block  
 ↳ facilitates recovery of at most 1 disk failure  
 ↳ write performance low due to the need to write all parity data to a single disk



→ RAID 5 : distributes parity blocks uniformly over all disks  
 ↳ several write requests simultaneously  
 ↳ read will have high parallelism  
 ↳ recovery of 1 disk failure

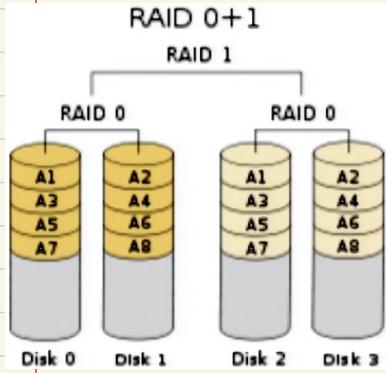


→ RAID 6 : two parity blocks distributed across all member disks  
 ↳ write performance poorer than RAID 5  
 ↳ recovery of 2 disk failures

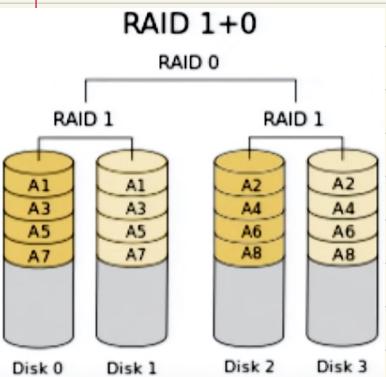
$$\begin{array}{r} 0100 \quad 0101 \\ 6 \quad 2 \\ \hline 0101 \quad 1000 \end{array} \quad \begin{array}{r} 0101 \quad 0000 \\ 6 \quad 6 \\ \hline 0101 \quad 0000 \end{array} \quad \begin{array}{r} 1000 \quad 0100 \\ 1101 \quad 0101 \\ \hline 1101 \quad 0000 \end{array}$$

$$\begin{array}{r} 0100 \quad 0100 \\ 0101 \quad 0101 \\ \hline 0001 \quad 0001 \end{array} \quad \begin{array}{r} 0001 \quad 0001 \\ 1101 \quad 0101 \\ \hline 1101 \quad 0100 \end{array}$$

→ Hybrid RAID : Nested RAID levels  
 ↳ combine two or more  
 ↳ example: RAID 50 layers the data stripping of RAID 0 on top of the distributed parity of RAID 5.

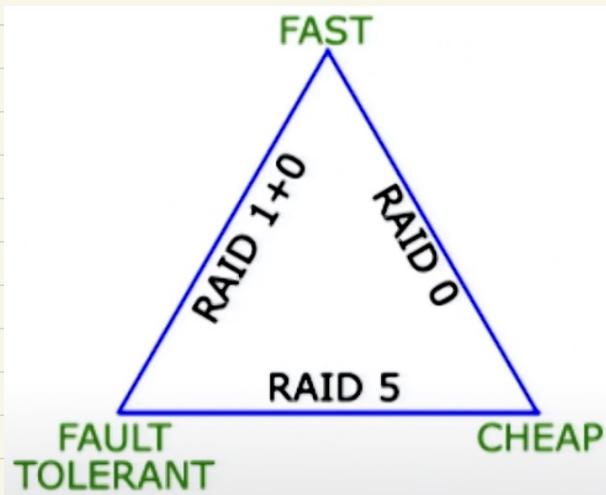


→ RAID 01 : mirror of stripes  
 ↳ both replication and sharing of data



→ RAID 10 : Stripe of mirroring  
 ↳ mirroring on striped data  
 ↳ better throughput and latency than all RAID levels except RAID 0.

→ Choice of RAID levels



→ What does RAID not do?

- doesn't equate 100% uptime
- " replace backups
- " protect against data corruption
- " allow to dynamically increase array size
- isn't the best option for virtualisation

# Comparison of RAID: Theoretical

PPD

| Level  | Description                                               | Min. <sup>[b]</sup> #<br>of drives | Space Efficiency             | Fault<br>Tolerance<br>(Drives) | Performance                |                                                            |
|--------|-----------------------------------------------------------|------------------------------------|------------------------------|--------------------------------|----------------------------|------------------------------------------------------------|
|        |                                                           |                                    |                              |                                | Read                       | Write                                                      |
|        |                                                           |                                    |                              |                                | (as factor of single disk) |                                                            |
| RAID 0 | Block-level striping without parity or mirroring          | 2                                  | 1                            | None                           | $n$                        | $n$                                                        |
| RAID 1 | Mirroring without parity or striping                      | 2                                  | $\frac{1}{n}$                | $n - 1$                        | $n^{[a]}$                  | $1^{[c]}$                                                  |
| RAID 2 | Bit-level striping with Hamming code for error correction | 3                                  | $1 - \frac{1}{n} \lg(n + 1)$ | One <sup>[d]</sup>             | Depends                    | Depends                                                    |
| RAID 3 | Byte-level striping with dedicated parity                 | 3                                  | $1 - \frac{1}{n}$            | One                            | $n - 1$                    | $n - 1^{[e]}$                                              |
| RAID 4 | Block-level striping with dedicated parity                | 3                                  | $1 - \frac{1}{n}$            | One                            | $n - 1$                    | $n - 1^{[e]}$                                              |
| RAID 5 | Block-level striping with distributed parity              | 3                                  | $1 - \frac{1}{n}$            | One                            | $n^{[e]}$                  | single sector: $\frac{1}{4}$<br>full stripe: $n - 1^{[e]}$ |
| RAID 6 | Block-level striping with double distributed parity       | 4                                  | $1 - \frac{2}{n}$            | Two                            | $n^{[e]}$                  | single sector: $\frac{1}{6}$<br>full stripe: $n - 2^{[e]}$ |

# Comparison of RAID: Practical

PPD

| Features             | RAID 0                                                                                | RAID 1                                          | RAID 5                                        | RAID 6                                                                                                     | RAID 10                                                  |
|----------------------|---------------------------------------------------------------------------------------|-------------------------------------------------|-----------------------------------------------|------------------------------------------------------------------------------------------------------------|----------------------------------------------------------|
| Minimum # of drives  | 2                                                                                     | 2                                               | 3                                             | 4                                                                                                          | 4                                                        |
| Fault tolerance      | None                                                                                  | Single-drive failure                            | Single-drive failure                          | Two-drive failure                                                                                          | Up to 1 disk failure in each sub-array                   |
| Read performance     | High                                                                                  | Medium                                          | Low                                           | Low                                                                                                        | High                                                     |
| Write Performance    | High                                                                                  | Medium                                          | Low                                           | Low                                                                                                        | Medium                                                   |
| Capacity utilization | 100%                                                                                  | 50%                                             | 67% – 94%                                     | 50% – 88%                                                                                                  | 50%                                                      |
| Typical applications | <i>High end workstations, data logging, real-time rendering, very transitory data</i> | <i>Operating systems, transaction databases</i> | <i>Data warehouse, web servers, archiving</i> | <i>Data archive, backup to disk, high availability solutions, servers with large capacity requirements</i> | <i>Fast databases, file servers, application servers</i> |



# WEEK 12

## → Overview of Query Processing

→ Steps: Parsing/Translator → Optimization → Evaluation  
 → Optimizer

↳ ex.  $\text{SELECT salary FROM instructor WHERE salary < 75000}$

→  $\text{O}_{\text{salary} < 75000} (\Pi_{\text{salary}}(\text{instructor}))$  ] equivalent plans  
 $\Pi_{\text{salary}} (\text{O}_{\text{salary} < 75000} (\text{instructor}))$  ]

↳ amongst all equivalent plans, choose the one with the lowest cost.

## → Measures of Query Cost

→ measured in terms of total elapsed time

↳ disk access is the primary focus

→ number of block transfers from disk and the number of seeks

→  $t_r$ : time to transfer one block

$t_s$ : time for one seek

→ cost of  $b$  block transfers plus  $S$  seeks =  $b \cdot t_r + S \cdot t_s$

## → Selection Operation

| A# | Algorithm                    | Cost                                       | Reason                                                                                                                                                                                                               |
|----|------------------------------|--------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| A1 | Linear Search                | $t_s + b \times t_r$                       | One initial seek plus $b$ block transfers                                                                                                                                                                            |
| A1 | Linear Search,<br>Eq. on Key | Average case<br>$t_s + (b_r/2) \times t_r$ | Since at most one record satisfies condition, scan can be terminated as soon as the required record is found. $b_r$ blocks transfers in worst case                                                                   |
| A2 | Prm. Index,<br>Eq. on Key    | $(h_i + 1) \times (t_r + t_s)$             | Index lookup traverses the height of the tree plus one I/O to fetch the record; each of these I/O operations requires a seek and a block transfer                                                                    |
| A3 | Prm. Index,<br>Eq. on Nonkey | $h_i \times (t_r + t_s) + b \times t_r$    | One seek for each level of the tree, one seek for the first block. Here all of $b$ are read. These blocks are leaf blocks assumed to be stored sequentially (for a primary index) and don't require additional seeks |
| A4 | Snd. Index,<br>Eq. on Key    | $(h_i + 1) \times (t_r + t_s)$             | This case is similar to primary index                                                                                                                                                                                |
| A4 | Snd. Index,<br>Eq. on Nonkey | $(h_i + n) \times (t_r + t_s)$             | Here, cost of index traversal is the same as for A3, but each record may be on a different block, requiring a seek per record. Cost is potentially very high if $n$ is large                                         |
| A5 | Prm. Index,<br>Comparison    | $h_i \times (t_r + t_s) + b \times t_r$    | Identical to the case of A3, equality on nonkey                                                                                                                                                                      |
| A6 | Snd. Index,<br>Comparison    | $(h_i + n) \times (t_r + t_s)$             | Identical to the case of A4, equality on nonkey                                                                                                                                                                      |

$t_r$  is time to transfer one block.  $t_s$  is time for one seek

$b$  denotes the number of blocks in the file

$b_r$  denotes the number of blocks containing records with the specified search key

$h_i$  denotes the height of the index.  $n$  is the number of records fetched

## → Conjunction

↳ A7 conjunctive selection using one index

A8 " " " composite index

A9 " " " by intersection of identifiers

## Disjunction

↳ A10 disjunctive selection by union of identifiers

↳ applicable if all conditions have available indices otherwise use linear scan

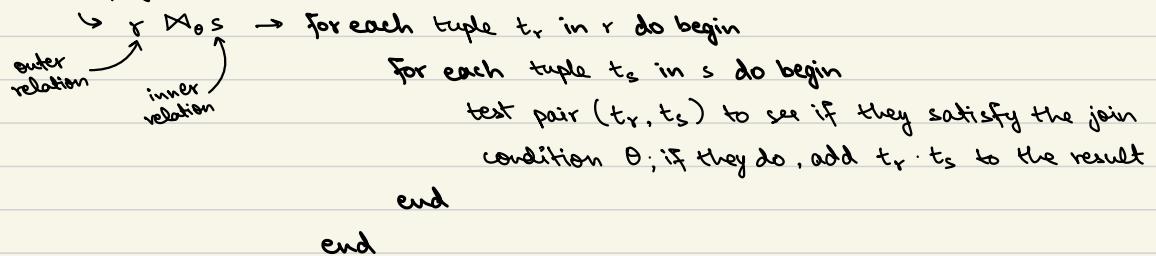
## → Sorting

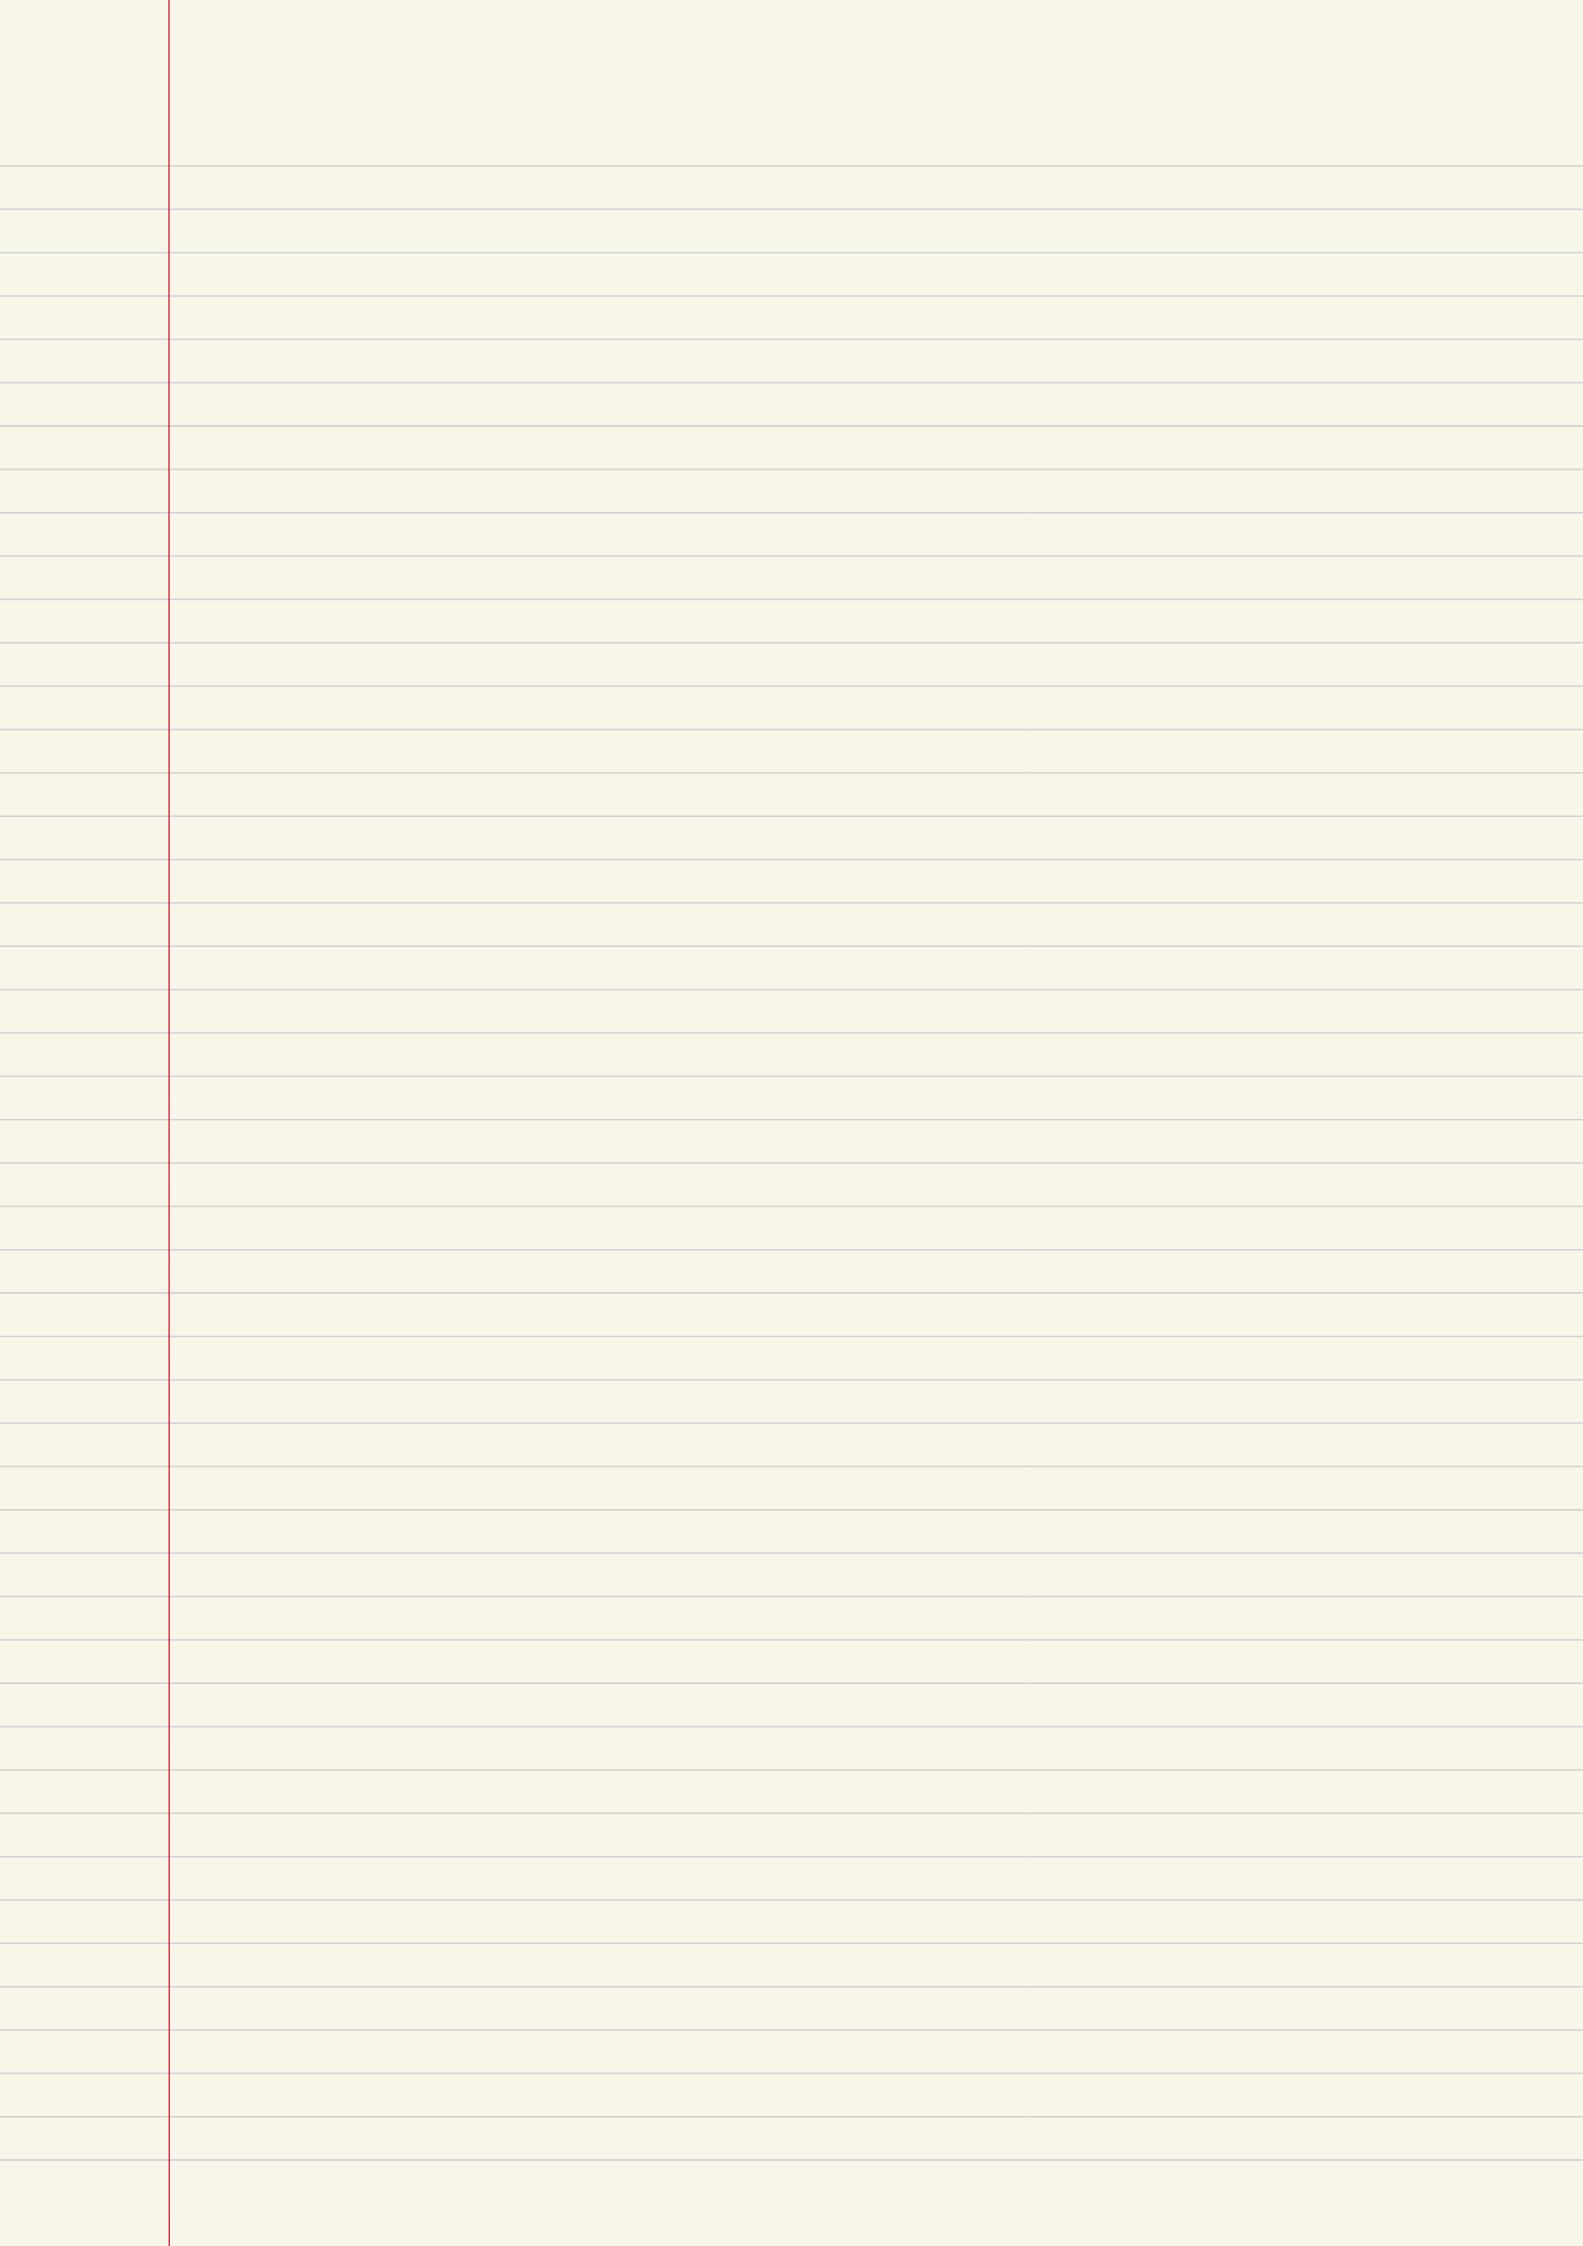
→ Quicksort : if relation fits in memory

External sort-merge : if relation is too big to fit in memory

## → Join

→ Nested-loop join





## WEEK 5 GA

$$\textcircled{1} \quad (P)^+ = PQCAB \quad (B)^+ = BPQCA \quad (C)^+ = C \\ (Q)^+ = QBPQA \quad (A)^+ = AQ \quad (AC)^+ = ACBPA$$

\textcircled{2}  $R(\overset{A}{\text{studInfo}}, \overset{B}{\text{courseId}}, \overset{C}{\text{name}}, \overset{D}{\text{state}})$

super keys  $\rightarrow A, AB, AC, AD$

$$\# = \sum_{i=0}^3 |C_i| = 1 + 3 + 3 + 1 = 8$$

$$\textcircled{4} \quad (PR)^+ = PRQST$$

\textcircled{5} any superkey must have 'A'.

$$(A)^+ = A \quad (AY)^+ = AYZBCX \quad (AB)^+ = ABCXYZ \quad (AZ)^+ = AZBCXY \\ (AC)^+ = ACXYZB \quad (XY)^+ = XYZBC \quad (AX)^+ = AXYZBC$$

\textcircled{6} if C is extraneous,  $F' = A \rightarrow C, AB \rightarrow D$

$$(AB)^+ \text{ w.r.t. } F' = ABDC$$

$$(AB)^+ \text{ w.r.t. } F = ABCD$$

if D is extraneous,  $F' = A \rightarrow C, AB \rightarrow C$

$$(AB)^+ \text{ w.r.t. } F' = ABC$$

$$(AB)^+ \text{ w.r.t. } F = ABCD$$

$$\textcircled{7} \quad (A)^+ = ABCDE$$

since A is not on the R.H.S. of any FD, A must be in every superkey



Since C is only being determined by A, drop A  $\rightarrow$  C.

— X — X — X —

$$A \rightarrow B$$

## WEEK 6 GA

$$2NF: A \rightarrow C$$

AND  
B non-prime

$$3NF: A \rightarrow SK$$

OR  
B  $\rightarrow$  prime

$$BCNF: A \rightarrow SK$$

$$\textcircled{1} \quad CK = AB \Rightarrow BCNF$$

$$CK = AB \Rightarrow 1NF$$

$$CK = AB \Rightarrow 2NF$$

$$CK = AB, DE \Rightarrow 3NF$$

$$\textcircled{2} \quad IPI(A, B, C, D) \quad FD = \{ A \rightarrow BCD, BC \rightarrow AD, D \rightarrow B \}$$

$$CK = A, BC, CD$$

$$2NF \checkmark$$

$$3NF \checkmark$$

$$BCNF \times \text{ because of } D \rightarrow B$$

$$\textcircled{3} \quad CK = T, PQ$$

$$2NF \checkmark \quad 3NF \times$$

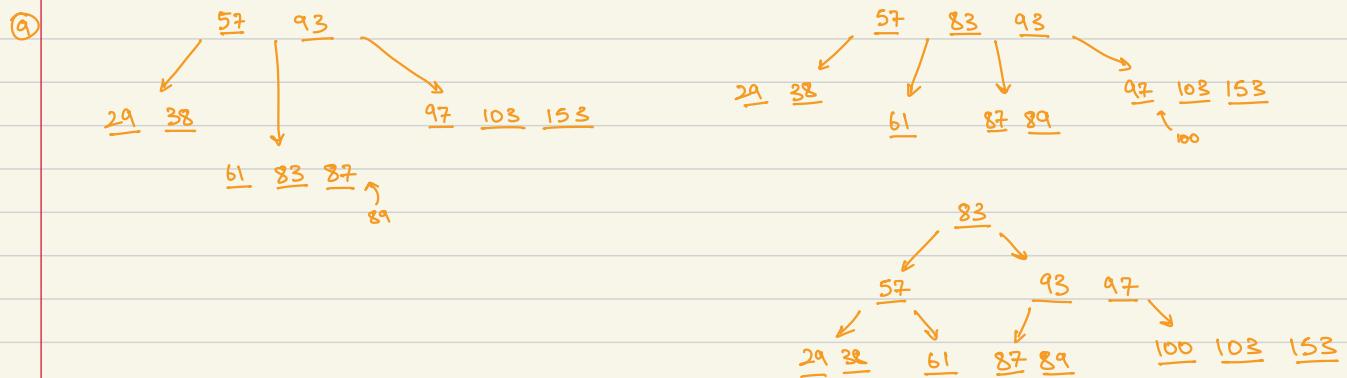
④  $R \rightarrow S$   
 $\uparrow$        $\downarrow$   
 $Q \leftarrow P$

$D_1 = (P, Q)$ ,  $FD = \{P \rightarrow Q\}$  BCNF ✓  
 $D_2 = (R, S)$ ,  $FD = \{R \rightarrow S\}$  BCNF ✓

⑤  $X(A, B, C, D, E, F)$      $CK = A, BC$      $2NF \checkmark$  BCNF ✓  
 $Y(G, H)$                    $CK = G$                    $2NF \checkmark$  BCNF ✓  
 $Z(G, A, I)$                    $CK = GA$                    $2NF \checkmark$  BCNF ✓  
 $FD = \{A \rightarrow BCDEF, BC \rightarrow ADEF,$   
 $G \rightarrow H,$   
 $GA \rightarrow I\}$

— X — X — X —

### WEEK 9 GA



— X — X — X —

### WEEK 10 GA

